

QUANTITATIVE_FINANCE_AND_RESEARCH

July 10, 2025

1 WEEK 1

I chose Apple, Microsoft, Amazon, Google and Tesla as the five large-cap US companies whose stocks I analyzed for my project. The link to the dataset I worked with is attached below. I did my project in the following steps:

1. Data Access
2. Data Cleaning
3. Data Transformation
4. Data Analysis

Link: <https://www.kaggle.com/datasets/borismarjanovic/price-volume-data-for-all-us-stocks-etfs>

1.1 DATA ACCESS

Here, I created a data frame comprising the data of stocks of my chosen companies, with ticker being the outer index and date being the inner index. Then I sorted the data in a descending manner, and converted the data in date column to DateTime format.

```
[98]: import pandas as pd
tickers = ['aapl', 'amzn', 'googl', 'msft', 'tsla']
dfs = []
for ticker in tickers:
    file_path = f"{ticker}.us.txt"
    try:
        df = pd.read_csv(file_path)
        df['ticker'] = ticker
        df['date'] = pd.to_datetime(df['Date'])
        df = df[['date', 'ticker', 'Open', 'High', 'Low', 'Close',
↪ 'Volume', 'OpenInt']]
        df.columns = ['Date', 'Ticker', 'Open', 'High', 'Low', 'Close',
↪ 'Volume', 'Open Interest']
        dfs.append(df)
    except Exception as e:
        print(f"Error: {e}")
combined_df = pd.concat(dfs)
multiindex_df = combined_df.set_index(['Ticker', 'Date']).sort_index(ascending=
↪ False)
```

1.2 DATA CLEANING

First, I checked for any missing rows across the tickers. Since there were none, I then looked for duplicates. I found only six duplicates outside the range of the years we needed to study, so I removed them. Next, I examined the columns for any that contained only one value, as these would be redundant. I also dropped those columns and checked for any columns with multiple data types that needed attention; there were none. Therefore, I filtered the data to include only the past 10 years, removed the other outliers, and proceeded with my analysis.

```
[99]: multiindex_df.shape
```

```
[99]: (26691, 6)
```

```
[100]: multiindex_df.isnull().sum()
```

```
[100]: Open          0
      High          0
      Low           0
      Close         0
      Volume        0
      Open Interest  0
      dtype: int64
```

```
[101]: duplicate_rows = multiindex_df[multiindex_df.duplicated()]
      if not duplicate_rows.empty:
          print("Duplicate rows found:")
          print(duplicate_rows)
      else:
          print("No duplicate rows found.")
```

Duplicate rows found:

		Open	High	Low	Close	Volume	Open Interest
Ticker Date							
msft	1986-09-16	0.07533	0.07533	0.07533	0.07533	6889952	0
	1986-07-07	0.08389	0.08389	0.07533	0.07533	18679429	0
	1986-07-03	0.08389	0.08389	0.08389	0.08389	18794263	0
	1986-05-15	0.08389	0.08389	0.08389	0.08389	5052632	0
	1986-04-24	0.07533	0.08389	0.07533	0.08389	82870827	0
	1986-04-09	0.07533	0.07533	0.07533	0.07533	16153115	0

```
[102]: multiindex_df.drop_duplicates(inplace = True)
```

```
[103]: for column in multiindex_df.columns:
      unique_values = multiindex_df[column].unique()
      print(f"Column '{column}': {len(unique_values)} unique values")
```

```
Column 'Open': 16297 unique values
Column 'High': 16188 unique values
```

```

Column 'Low': 16280 unique values
Column 'Close': 16651 unique values
Column 'Volume': 26120 unique values
Column 'Open Interest': 1 unique values

```

```
[104]: multiindex_df = multiindex_df.drop(columns=['Open Interest'])
```

```
[105]: def check_datatypes(df):
        print("Data types of DataFrame columns:")
        print(df.dtypes)
        if isinstance(df.index, pd.MultiIndex):
            print("\nData types of MultiIndex levels:")
            for i, level_name in enumerate(df.index.names):
                print(f"Level '{level_name}': {df.index.get_level_values(i).dtype}")
        check_datatypes(multiindex_df)
```

Data types of DataFrame columns:

```

Open      float64
High      float64
Low       float64
Close     float64
Volume    int64
dtype: object

```

Data types of MultiIndex levels:

```

Level 'Ticker': object
Level 'Date': datetime64[ns]

```

```
[163]: import datetime
        filtered_dfs = []
        for ticker in tickers:
            try:
                df_ticker = multiindex_df.loc[ticker]
                latest_date = df_ticker.index.max()
                date_10_years_ago = latest_date - pd.DateOffset(years=10)
                df_filtered_ticker = df_ticker[(df_ticker.index <= latest_date) &
↪ (df_ticker.index >= date_10_years_ago)]
                filtered_dfs.append(df_filtered_ticker)
            except KeyError:
                print(f"Ticker '{ticker}' not found in the DataFrame index.")
            except Exception as e:
                print(f"An error occurred for ticker '{ticker}': {e}")
        filtered_df = pd.concat(filtered_dfs, keys=tickers, names=['Ticker', 'Date'])
```

```
[107]: def remove_outliers_iqr(df, column):
        Q1 = df[column].quantile(0.25)
        Q3 = df[column].quantile(0.75)
```

```

IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
df_filtered = df[(df[column] >= lower_bound) & (df[column] <= upper_bound)]
return df_filtered
filtered_df = remove_outliers_iqr(filtered_df, 'Close')
filtered_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
MultiIndex: 10980 entries, ('aapl', Timestamp('2017-11-10 00:00:00')) to
('tsla', Timestamp('2010-06-28 00:00:00'))
Data columns (total 5 columns):
#   Column   Non-Null Count  Dtype
---  -
0   Open     10980 non-null  float64
1   High     10980 non-null  float64
2   Low      10980 non-null  float64
3   Close    10980 non-null  float64
4   Volume   10980 non-null  int64
dtypes: float64(4), int64(1)
memory usage: 545.6+ KB

```

1.3 DATA TRANSFORMATION

I added some new important features to my given dataset, removed some unimportant ones and methodologically answered some fundamental questions regarding the given dataset.

For each stock, I added the following columns: 1. Daily Return: % change in closing price 2. 7-day Moving Average of closing price 3. 30-day Moving Average of closing price 4. Rolling Volatility (30d): Standard deviation of returns over the last 30 days

```

[108]: filtered_df = filtered_df.copy()
filtered_df_sorted = filtered_df.sort_index(level=['Ticker', 'Date'],
        ↪ascending=[True, True])
filtered_df_sorted['Daily Return'] = filtered_df_sorted.
        ↪groupby('Ticker')['Close'].pct_change() * 100

```

```

[109]: filtered_df_sorted['7-Day Moving Average'] = (
        filtered_df_sorted.groupby('Ticker')['Close']
        .transform(lambda x: x.rolling(window=7).mean())
    )
filtered_df_sorted['30-Day Moving Average'] = (
        filtered_df_sorted.groupby('Ticker')['Close']
        .transform(lambda x: x.rolling(window=30).mean())
    )
filtered_df_sorted['Rolling Volatility (30d)'] = (
        filtered_df_sorted.groupby('Ticker')['Daily Return']

```

```
.transform(lambda x: x.rolling(window=30).std())
)
```

Then, I answered the following questions:

1. Which stock had the highest average return over the 10 years?
2. Which stock had the most volatile month, and when?

```
[110]: average_returns = filtered_df_sorted.groupby('Ticker')['Daily Return'].mean()
highest_avg_return_stock = average_returns.idxmax()
print(f"The stock with the highest average return over the 10-year period is:␣
      ↳{highest_avg_return_stock}")
print(f"Average returns:␣n{average_returns}")
```

The stock with the highest average return over the 10-year period is: tsla

Average returns:

Ticker

aapl	0.106375
amzn	0.134835
googl	0.056440
msft	0.058938
tsla	0.210083

Name: Daily Return, dtype: float64

```
[111]: volatility_unstacked = filtered_df_sorted['Rolling Volatility (30d)'].
      ↳unstack(level='Ticker')
monthly_volatility = volatility_unstacked.resample('M').mean()
average_monthly_volatility_across_tickers = monthly_volatility.mean(axis=1)
most_volatile_month_date = average_monthly_volatility_across_tickers.idxmax()
volatility_in_most_volatile_month = monthly_volatility.
      ↳loc[most_volatile_month_date]
most_volatile_stock_in_month = volatility_in_most_volatile_month.dropna().
      ↳idxmax()
print(f"␣nThe stock with the most volatile month was␣
      ↳{most_volatile_stock_in_month}.")
print(f"This occurred in {most_volatile_month_date.strftime('%B %Y')}.")
print(f"The average rolling volatility across all stocks in that month was:␣
      ↳{average_monthly_volatility_across_tickers.max():.2f}")
print(f"The rolling volatility of {most_volatile_stock_in_month} in that month␣
      ↳was: {volatility_in_most_volatile_month.max():.2f}")
```

The stock with the most volatile month was amzn.

This occurred in November 2008.

The average rolling volatility across all stocks in that month was: 5.72

The rolling volatility of amzn in that month was: 6.45

Daily returns are calculated assuming markets are liquid, meaning there are no gaps between

closing prices. Rolling volatility, a key metric in our analysis, is based on the assumption that returns are typically distributed over 30-day periods. This period is crucial in understanding the market dynamics and reassures the thoroughness of our analysis. Please note that this analysis does not consider factors such as dividend payments, stock splits, or after-hours trading, which could impact return calculations.

After performing an exploratory analysis of the given data, I found out that:

1. The stock with the highest average return over the 10 years was TSLA.
2. The stock with the most volatile month was AMZN in November 2008.

Then I proceeded to add some more features to my column, convert the categorical values, check for skewness of my dataset and restructured my dataset.

```
[112]: skewness = filtered_df_sorted['Close'].skew()
print(f"Skewness of the 'Close' column: {skewness}")
if skewness > 0:
    print("The 'Close' column is right-skewed (positive skewness).")
elif skewness < 0:
    print("The 'Close' column is left-skewed (negative skewness).")
else:
    print("The 'Close' column is not skewed (or very close to symmetrical).")
```

Skewness of the 'Close' column: 1.2505667180241913
The 'Close' column is right-skewed (positive skewness).

```
[113]: import numpy as np
filtered_df_sorted['Close_Log'] = np.log(filtered_df_sorted['Close'])
skewness = filtered_df_sorted['Close_Log'].skew()
print(f"Skewness of the 'Close_Log' column: {skewness}")
if skewness > 0:
    print("The 'Close_Log' column is right-skewed (positive skewness).")
elif skewness < 0:
    print("The 'Close_Log' column is left-skewed (negative skewness).")
else:
    print("The 'Close_Log' column is not skewed (or very close to symmetrical).")
```

Skewness of the 'Close_Log' column: -0.031109674242157298
The 'Close_Log' column is left-skewed (negative skewness).

```
[114]: filtered_df_sorted['Year'] = filtered_df_sorted.index.get_level_values('Date').
    ↪year
filtered_df_sorted['Month'] = filtered_df_sorted.index.get_level_values('Date').
    ↪month
filtered_df_sorted['Day'] = filtered_df_sorted.index.get_level_values('Date').
    ↪dayofyear
filtered_df_sorted['Day_of_Week'] = filtered_df_sorted.index.
    ↪get_level_values('Date').dayofweek
```

```

filtered_df_sorted['Number_of_Days'] = (filtered_df_sorted.index.
    ↪get_level_values('Date') - filtered_df_sorted.index.get_level_values('Date')).
    ↪min()).days

```

```

[115]: filtered_df_sorted["Price Change"] = filtered_df_sorted["Close"] -
    ↪filtered_df_sorted["Open"]
filtered_df_sorted["High-Low Spread"] = filtered_df_sorted["High"] -
    ↪filtered_df_sorted["Low"]
filtered_df_sorted["Intraday Volatility %"] = (filtered_df_sorted["High"] -
    ↪filtered_df_sorted["Low"]) / filtered_df_sorted["Open"] * 100
filtered_df_sorted["Lag_1_Daily_Return"] = filtered_df_sorted["Daily_Return"].
    ↪shift(1)
def compute_rsi(series, period=14):
    delta = series.diff()
    gain = delta.clip(lower=0)
    loss = -1 * delta.clip(upper=0)
    avg_gain = gain.rolling(window=period).mean()
    avg_loss = loss.rolling(window=period).mean()
    rs = avg_gain / avg_loss
    return 100 - (100 / (1 + rs))
filtered_df_sorted["RSI_14"] = compute_rsi(filtered_df_sorted["Close"], 14)
exp1 = filtered_df_sorted["Close"].ewm(span=12, adjust=False).mean()
exp2 = filtered_df_sorted["Close"].ewm(span=26, adjust=False).mean()
filtered_df_sorted["MACD"] = exp1 - exp2
filtered_df_sorted["Volume_Change"] = filtered_df_sorted["Volume"].pct_change()
filtered_df_sorted["Volume_Price_Ratio"] = filtered_df_sorted["Volume"] /
    ↪filtered_df_sorted["Close"]

```

```

[116]: for col in filtered_df_sorted.columns:
    print(f"Number of unique values in {col} are: {filtered_df_sorted[col].
    ↪nunique()}")

```

```

Number of unique values in Open are: 9670
Number of unique values in High are: 9581
Number of unique values in Low are: 9667
Number of unique values in Close are: 9824
Number of unique values in Volume are: 10970
Number of unique values in Daily_Return are: 10912
Number of unique values in 7-Day Moving Average are: 10852
Number of unique values in 30-Day Moving Average are: 10817
Number of unique values in Rolling Volatility (30d) are: 10830
Number of unique values in Close_Log are: 9824
Number of unique values in Year are: 11
Number of unique values in Month are: 12
Number of unique values in Day are: 365
Number of unique values in Day_of_Week are: 5
Number of unique values in Number_of_Days are: 2519

```


Number of unique values in Price Change are: 5798
 Number of unique values in High-Low Spread are: 5226
 Number of unique values in Intraday Volatility % are: 10973
 Number of unique values in Lag_1_Daily_Return are: 10912
 Number of unique values in RSI_14 are: 10941
 Number of unique values in MACD are: 10980
 Number of unique values in Volume_Change are: 10971
 Number of unique values in Volume_Price_Ratio are: 10976

```
[117]: cols = [col for col in filtered_df_sorted.columns if (col != 'Close' and col != 'Close_Log')] + ['Close'] + ['Close_Log']
final_df = filtered_df_sorted[cols]
final_df.to_csv('stock_data.csv')
```


1.4 EXPLORATORY DATA ANALYSIS

While this analysis focuses on quantitative metrics, visualisations (e.g., moving average trends or volatility time series) could further illustrate patterns like Tesla's consistent outperformance or Amazon's volatility spikes during the 2008 financial crisis. Thus, I added some visual representations to explain the above patterns in this section, to make my explanations and assessments of the given data clear, and to ease the reader's understanding of the data. I also analysed other stuff, such as the correlation matrix and skewness of the close column.


```
[118]: import matplotlib.pyplot as plt
tsla_data = filtered_df_sorted.loc['tsla']
tsla_data.plot(y=['Close', '30-Day Moving Average'], figsize=(12, 6))
plt.title('Tesla: Price vs. 30-Day Moving Average')
plt.ylabel('Price')
plt.show()
```




```
[119]: amzn_monthly_volatility = monthly_volatility['amzn']
plt.figure(figsize=(15, 7))
amzn_monthly_volatility.plot(kind='bar')
plt.title('Amazon Monthly Rolling Volatility (30d)')
plt.xlabel('Date')
plt.ylabel('Volatility')
labels = [date.strftime('%Y') if date.month == 1 else '' for date in_
↪ amzn_monthly_volatility.index]
plt.xticks(range(len(amzn_monthly_volatility.index)), labels, rotation=0)
plt.tight_layout()
plt.show()
```




```
[120]: import seaborn as sns
sns.histplot(filtered_df_sorted['Close'], kde=True)
plt.title("Distribution of Closing Prices")
plt.show()
from scipy.stats import skew
print("Skewness of Close:", skew(filtered_df_sorted['Close']))
```



Skewness of Close: 1.2503958690928683

```
[121]: plt.figure(figsize=(15, 7))
for ticker in tickers:
    try:
        df_ticker = filtered_df_sorted.loc[ticker].sort_index(ascending=True)
        plt.plot(df_ticker.index.get_level_values('Date'),
df_ticker['Close_Log'], label=ticker)
    except KeyError:
        print(f"Data for ticker '{ticker}' not found.")

plt.title('Log of Closing Price Over Time for Selected Stocks')
plt.xlabel('Date')
plt.ylabel('Log(Close Price)')
plt.legend()
plt.grid(True)
plt.show()
```



2 WEEK 2

I chose the Amazon stock data, split it into training and testing datasets and implemented three models: Linear Regression, Random Forest and ARIMA. I picked Random Forest as the best model based on the Mean Absolute Error (MAE) evaluation metrics and direction prediction accuracy. I backtested it on the data: buy if the predicted price exceeds the current price; otherwise, sell. I also calculated the hypothetical profit/loss over the test period.

2.1 MODEL IMPLEMENTATION AND EVALUATION

I implemented and evaluated ARIMA, Linear Regression and Random Forest Models based on Mean Absolute Error (MAE) and direction prediction accuracy as the two evaluation metrics.

2.1.1 DATA SPLITTING

I chose AMZN as the stock to test my model. I split the data into the training (first 80%) and testing (last 20%) sets, ensuring the split respects the time order.

```
[122]: features_random_forest = [
    'Open', 'High', 'Low',
    '7-Day Moving Average', '30-Day Moving Average',
    'High-Low Spread', 'MACD',
    'Rolling Volatility (30d)', 'Intraday Volatility %',
    'Volume', 'Volume_Price_Ratio',
    'Number_of_Days', 'Year'
]
```

```

features_linear_regressor = ['Open', '7-Day Moving Average', 'High-Low Spread',
    ↪ 'MACD', 'Volume', 'Volume_Price_Ratio']
amzn_df = final_df.loc['amzn'].copy()
amzn_df = amzn_df[amzn_df['Close'] > 0]
amzn_df.dropna(subset=features_random_forest, inplace=True)
amzn_df['Close_Log'] = np.log(amzn_df['Close'])
amzn_df.dropna(subset=['Close_Log'], inplace=True)
amzn_df.reset_index(inplace=True)
train_size = int(len(amzn_df) * 0.8)
train_df = amzn_df.iloc[:train_size]
test_df = amzn_df.iloc[train_size:]

```

2.1.2 MODEL IMPLEMENTATION AND EVALUATION

I implemented three models, linear regression, ARIMA, and random forest, and evaluated them. Evaluation would be done using Mean Absolute Error (MAE) and accuracy for direction prediction (whether the price goes up or down).

LINEAR REGRESSION

```

[123]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train = train_df[features_linear_regressor]
y_train = train_df["Close_Log"]
X_test = test_df[features_linear_regressor]
y_test = test_df["Close_Log"]
X_train_scaled = scaler.fit_transform(train_df[features_linear_regressor])
X_test_scaled = scaler.transform(test_df[features_linear_regressor])
y_scaler = StandardScaler()
y_train_scaled = y_scaler.fit_transform(train_df[['Close_Log']])
y_test_scaled = y_scaler.transform(test_df[['Close_Log']])
X_train_scaled = pd.DataFrame(X_train_scaled, columns=features_linear_regressor)
X_test_scaled = pd.DataFrame(X_test_scaled, columns=features_linear_regressor)
y_train_scaled = pd.DataFrame(y_train_scaled, columns=['Close_Log'])
y_test_scaled = pd.DataFrame(y_test_scaled, columns=['Close_Log'])

```

```

[124]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error
def create_lagged_features(df, lag=5):
    df_lagged = df.copy()
    for i in range(1, lag + 1):
        df_lagged[f'Close_Lag_{i}'] = df_lagged['Close_Log'].shift(i)
    df_lagged.dropna(inplace=True)
    return df_lagged
original_test_index = test_df.index.copy()
test_df_reset = test_df.reset_index(drop=True)
train_df_lagged = create_lagged_features(pd.concat([X_train_scaled,
    ↪ y_train_scaled], axis=1), lag=5)

```

```

test_df_lagged = create_lagged_features(pd.concat([X_test_scaled,
↳ y_test_scaled], axis=1), lag=5)
lagged_features = [f'Close_Lag_{i}' for i in range(1, 6)]
X_train_lagged = train_df_lagged[lagged_features]
y_train_lagged = train_df_lagged['Close_Log']
X_test_lagged = test_df_lagged[lagged_features]
y_test_lagged = test_df_lagged['Close_Log']
lr_model = LinearRegression()
lr_model.fit(X_train_lagged, y_train_lagged)
y_pred_scaled = lr_model.predict(X_test_lagged)
y_pred = y_scaler.inverse_transform(y_pred_scaled.reshape(-1, 1))
y_test_actual = y_scaler.inverse_transform(y_test_lagged.values.reshape(-1, 1))
mae_lr = mean_absolute_error(y_test_actual, y_pred)
print(f"Linear Regression MAE (using past 5 days' scaled close): {mae_lr:.4f}")
test_df_aligned = test_df_reset.loc[test_df_lagged.index]
actual_prev_close = test_df_aligned['Close_Log'].shift(1).dropna()
predicted_direction = np.sign(y_pred.flatten()[1:] - actual_prev_close)
actual_direction = np.sign(test_df_aligned['Close_Log'].values[1:] -
↳ actual_prev_close)
correct_predictions = np.sum(predicted_direction == actual_direction)
accuracy_lr = correct_predictions / len(actual_direction)
print(f"Linear Regression Direction Prediction Accuracy (using past 5 days'
↳ scaled close): {accuracy_lr:.4f}")

```

Linear Regression MAE (using past 5 days' scaled close): 0.0140

Linear Regression Direction Prediction Accuracy (using past 5 days' scaled close): 0.5499

RANDOM FOREST

```

[125]: from sklearn.ensemble import RandomForestRegressor
import random
import os
seed = 42
os.environ['PYTHONHASHSEED'] = str(seed)
np.random.seed(seed)
random.seed(seed)
target = 'Close_Log'
X_train_rf = train_df[features_random_forest].dropna()
y_train_rf = train_df.loc[X_train_rf.index, target]
X_test_rf = test_df[features_random_forest].dropna()
y_test_rf = test_df.loc[X_test_rf.index, target]
rf_model = RandomForestRegressor(
    n_estimators=150,
    max_features=0.8,
    min_samples_leaf=5,
    max_depth=12,
    min_samples_split=10,

```

```

    random_state=seed,
    n_jobs=1,
    bootstrap=True
)
rf_model.fit(X_train_rf, y_train_rf)
y_pred_rf = rf_model.predict(X_test_rf)
mae_rf = mean_absolute_error(y_test_rf, y_pred_rf)
print(f"Random Forest MAE: {mae_rf:.4f}")
actual_prev_close_rf = test_df.loc[y_test_rf.index].shift(1)[target].dropna()
aligned_y_pred_rf = pd.Series(y_pred_rf, index=y_test_rf.index).
    ↪loc[actual_prev_close_rf.index]
aligned_y_test_rf = y_test_rf.loc[actual_prev_close_rf.index]
predicted_direction_rf = np.sign(aligned_y_pred_rf - actual_prev_close_rf)
actual_direction_rf = np.sign(aligned_y_test_rf - actual_prev_close_rf)
non_zero_change_indices = actual_direction_rf != 0
correct_predictions_rf = np.sum(predicted_direction_rf[non_zero_change_indices]
    ↪== actual_direction_rf[non_zero_change_indices])
accuracy_rf = correct_predictions_rf / np.sum(non_zero_change_indices) if np.
    ↪sum(non_zero_change_indices) > 0 else 0
print(f"Random Forest Direction Prediction Accuracy: {accuracy_rf:.4f}")

```

Random Forest MAE: 0.1660

Random Forest Direction Prediction Accuracy: 0.6106

ARIMA

I will forecast the closing price using an ARIMA model.

```

[126]: import pmdarima as pm
model = pm.auto_arima(train_df['Close_Log'],
                      seasonal=True,
                      m=12,
                      trace=True,
                      error_action='ignore',
                      suppress_warnings=True)
n_periods = len(test_df)
forecast, conf_int = model.predict(n_periods=n_periods, return_conf_int=True)
test_df = test_df.copy()
test_df['Forecast_Close_Log'] = forecast
test_df['Forecast_Close'] = np.exp(test_df['Forecast_Close_Log'])
test_df['Actual_Close'] = np.exp(test_df['Close_Log'])
test_df.replace([np.inf, -np.inf], np.nan, inplace=True)
test_df.dropna(subset=['Forecast_Close', 'Actual_Close'], inplace=True)
mae_ar = mean_absolute_error(test_df['Actual_Close'], test_df['Forecast_Close'])
print(f"MAE: {mae_ar:.4f}")
test_df['Actual_Direction'] = test_df['Actual_Close'].diff().apply(lambda x: 1
    ↪if x > 0 else 0)

```

```

test_df['Predicted_Direction'] = test_df['Forecast_Close'].diff().apply(lambda x: 1 if x > 0 else 0)
test_df.dropna(subset=['Actual_Direction', 'Predicted_Direction'], inplace=True)
accuracy_ar = np.mean(test_df['Actual_Direction'] == test_df['Predicted_Direction'])
print(f"Directional Accuracy: {accuracy_ar * 100:.4f}%")
plt.figure(figsize=(12, 6))
plt.plot(train_df.index, np.exp(train_df['Close_Log']), label='Training Data')
plt.plot(test_df.index, test_df['Actual_Close'], label='Actual Close (Test)', color='blue')
plt.plot(test_df.index, test_df['Forecast_Close'], label='Forecasted Close', color='red')
plt.fill_between(test_df.index,
                 np.exp(conf_int[:, 0]),
                 np.exp(conf_int[:, 1]),
                 color='pink', alpha=0.3, label='Confidence Interval')
plt.legend()
plt.title('ARIMA Forecast vs Actual')
plt.xlabel('Date')
plt.ylabel('Stock Price')
plt.grid(True)
plt.tight_layout()
plt.show()

```

Performing stepwise search to minimize aic

```

ARIMA(2,1,2)(1,0,1)[12] intercept : AIC=-7339.187, Time=2.46 sec
ARIMA(0,1,0)(0,0,0)[12] intercept : AIC=-7340.723, Time=0.28 sec
ARIMA(1,1,0)(1,0,0)[12] intercept : AIC=-7338.411, Time=0.90 sec
ARIMA(0,1,1)(0,0,1)[12] intercept : AIC=-7338.680, Time=4.17 sec
ARIMA(0,1,0)(0,0,0)[12] intercept : AIC=-7341.449, Time=0.14 sec
ARIMA(0,1,0)(1,0,0)[12] intercept : AIC=-7338.728, Time=0.38 sec
ARIMA(0,1,0)(0,0,1)[12] intercept : AIC=-7338.703, Time=0.75 sec
ARIMA(0,1,0)(1,0,1)[12] intercept : AIC=-7336.709, Time=1.05 sec
ARIMA(1,1,0)(0,0,0)[12] intercept : AIC=-7340.424, Time=0.14 sec
ARIMA(0,1,1)(0,0,0)[12] intercept : AIC=-7340.680, Time=0.57 sec
ARIMA(1,1,1)(0,0,0)[12] intercept : AIC=-7339.648, Time=0.57 sec


```

Best model: ARIMA(0,1,0)(0,0,0)[12]

Total fit time: 11.435 seconds

MAE: 138.7132

Directional Accuracy: 46.7626%



2.2 MODEL COMPARISON

Based on the previous data, I compared the models based on MAE and identify the best-performing one. Then I saved it for backtesting.

```
[127]: linear_regression_mae = mae_lr
linear_regression_accuracy = accuracy_lr
random_forest_mae = mae_rf
random_forest_accuracy = accuracy_rf
arima_mae = mae_ar
arima_accuracy = accuracy_ar
print(f"Linear Regression MAE: {linear_regression_mae:.4f}")
print(f"Linear Regression Direction Accuracy: {linear_regression_accuracy:.4f}")
print(f"Random Forest MAE: {random_forest_mae:.4f}")
print(f"Random Forest Direction Accuracy: {random_forest_accuracy:.4f}")
print(f"ARIMA MAE: {arima_mae:.4f}")
print(f"ARIMA Directional Accuracy: {arima_accuracy:.4f}")
best_mae_model = min([
    ('Linear Regression', linear_regression_mae),
    ('Random Forest', random_forest_mae),
    ('ARIMA', arima_mae)
], key=lambda x: x[1])
print(f"\nModel with the lowest MAE: {best_mae_model[0]} (MAE: {best_mae_model[1]:.4f})")
best_accuracy_model = max([
    ('Linear Regression', linear_regression_accuracy),
    ('Random Forest', random_forest_accuracy),
    ('ARIMA', arima_accuracy)
], key=lambda x: x[1])
```



```

print(f"Model with the highest Direction Accuracy: {best_accuracy_model[0]}_
↳(Accuracy: {best_accuracy_model[1]:.4f})")
best_overall_model_name = best_accuracy_model[0]
print(f"\nBased on accuracy of the predicted direction, the best performing_
↳model is: {best_overall_model_name}")

```

Linear Regression MAE: 0.0140
 Linear Regression Direction Accuracy: 0.5499
 Random Forest MAE: 0.1660
 Random Forest Direction Accuracy: 0.6106
 ARIMA MAE: 138.7132
 ARIMA Directional Accuracy: 0.4676

Model with the lowest MAE: Linear Regression (MAE: 0.0140)
 Model with the highest Direction Accuracy: Random Forest (Accuracy: 0.6106)

Based on accuracy of the predicted direction, the best performing model is:
 Random Forest

2.3 BACKTESTING THE BEST MODEL

Using the best model, I generated buy/sell signals; Buy if the predicted price > current price, sell otherwise. I also calculated the hypothetical profit/loss over the test period.

```


[128]: test_df_backtest = test_df.loc[X_test_rf.index].copy()
test_df_backtest['Predicted_Close_Log_RF'] = y_pred_rf
test_df_backtest['Predicted_Close_RF'] = np.
↳exp(test_df_backtest['Predicted_Close_Log_RF'])
test_df_backtest['Previous_Close'] = test_df_backtest['Close'].shift(1)
test_df_backtest.dropna(subset=['Previous_Close'], inplace=True)
test_df_backtest['Signal'] = np.where(test_df_backtest['Predicted_Close_RF'] >_
↳test_df_backtest['Previous_Close'], 1, -1)
test_df_backtest['Strategy_Return'] = test_df_backtest['Signal'] *_
↳(test_df_backtest['Close'] - test_df_backtest['Previous_Close'])
test_df_backtest['Cumulative_Profit'] = test_df_backtest['Strategy_Return'].
↳cumsum()
test_df_backtest['Buy_Hold_Profit'] = (test_df_backtest['Close'] -_
↳test_df_backtest['Previous_Close']).cumsum()
print("Final Strategy Profit:", test_df_backtest['Cumulative_Profit'].iloc[-1])
print("Final Buy & Hold Profit:", test_df_backtest['Buy_Hold_Profit'].iloc[-1])
plt.figure(figsize=(12, 6))
plt.plot(test_df_backtest.index, test_df_backtest['Cumulative_Profit'],_
↳label='Strategy Profit')
plt.plot(test_df_backtest.index, test_df_backtest['Buy_Hold_Profit'],_
↳label='Buy & Hold Profit', linestyle='--')
plt.title('Backtesting Random Forest Strategy vs Buy & Hold')
plt.xlabel('Date')

```

```
plt.ylabel('Cumulative Profit (Assuming 1 Share)')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```

Final Strategy Profit: 455.89000000000002

Final Buy & Hold Profit: 339.72999999999996



3 WEEK 3

I developed more advanced features for multivariate modelling, including polynomial and exogenous variables. Then, I implemented several models: SARIMAX, an Ensemble Model, and a Temporal Fusion Transformer (TFT). Using the same evaluation metrics, I compared these models to those developed in Week 2. After that, I selected the best-performing model and backtested it on the dataset using the same strategy. I incorporated two additional components to manage risk: the stop-loss ratio and the Sharpe ratio. Finally, I calculated the expected returns and the covariance matrix and conducted Mean-Variance Optimization (MVO). The goal was to maximize the Sharpe ratio to determine the optimal asset weights for Apple and Amazon stocks.

3.1 ADVANCE FEATURE ENGINEERING

In this section, I decomposed the time series to extract weakly and monthly components, engineered polynomial features like SMA², RSI x Volume, etc.. and included some exogenous variables for multivariate modelling.

```
[129]: def load_yahoo_index(path, rename_map):
df = pd.read_csv(path, skiprows=1)
df = df[df['Ticker'] != 'Date']
```

```

df.rename(columns=rename_map, inplace=True)
df['Date'] = pd.to_datetime(df['Date'], format='%Y-%m-%d').dt.normalize()
return df
sp500 = load_yahoo_index('sp500_data.csv', {
    'Ticker': 'Date',
    '^GSPC': 'Open',
    '^GSPC.1': 'High',
    '^GSPC.2': 'Low',
    '^GSPC.3': 'Close',
    '^GSPC.4': 'Volume'
})
sp500.rename(columns={'Close': 'SP500_Close'}, inplace=True)
sp500.set_index('Date', inplace=True)
vix = load_yahoo_index('vix_data.csv', {
    'Ticker': 'Date',
    '^VIX': 'Open',
    '^VIX.1': 'High',
    '^VIX.2': 'Low',
    '^VIX.3': 'Close',
    '^VIX.4': 'Volume'
})
vix.rename(columns={'Close': 'VIX'}, inplace=True)
vix.set_index('Date', inplace=True)
from statsmodels.tsa.seasonal import STL
if 'Date' not in filtered_df_sorted.columns:
    filtered_df_sorted = filtered_df_sorted.reset_index()
for col in ['level_0', 'level_1', 'index']:
    if col in filtered_df_sorted.columns:
        filtered_df_sorted.drop(columns=[col], inplace=True)
filtered_df_sorted['Date'] = pd.to_datetime(filtered_df_sorted['Date']).dt.
    ↪normalize()
def add_decomposition_features_grouped(df, column='Close'):
    dfs = []
    if 'Date' not in df.columns:
        df = df.reset_index()
    for ticker, group in df.groupby('Ticker'):
        if 'Date' not in group.columns:
            group = group.reset_index()
        group = group.set_index('Date').sort_index()
        try:
            stl = STL(group[column], period=7, robust=True).fit()
            group[f'{column}_trend'] = stl.trend
            group[f'{column}_seasonal'] = stl.seasonal
            group[f'{column}_resid'] = stl.resid
        except Exception as e:
            print(f"STL failed for {ticker}: {e}")
            group[f'{column}_trend'] = None

```

```

        group[f'{column}_seasonal'] = None
        group[f'{column}_resid'] = None
        dfs.append(group.reset_index())
    return pd.concat(dfs, ignore_index=True)

def add_exogenous_features(df):
    market_df = sp500[['SP500_Close']].join(vix[['VIX']], how='inner')
    df = df.copy()
    if 'Date' not in df.columns:
        df = df.reset_index()
    df['Date'] = pd.to_datetime(df['Date']).dt.normalize()
    df = df.set_index('Date')
    df = df.merge(market_df, left_index=True, right_index=True, how='left')
    df = df.reset_index()
    return df

def add_lag_features(df):
    for lag in [1, 3, 5, 7]:
        df[f'Return_lag_{lag}'] = df.groupby('Ticker')['Daily Return'].
↳shift(lag)
        df[f'Volume_lag_{lag}'] = df.groupby('Ticker')['Volume'].shift(lag)
    return df

def add_difference_features(df):
    df['Volatility_diff'] = df.groupby('Ticker')['Rolling Volatility (30d)'].
↳diff()
    df['Volume_diff'] = df.groupby('Ticker')['Volume'].diff()
    return df

def full_feature_engineering(df):
    df = add_decomposition_features_grouped(df)
    df = add_exogenous_features(df)
    df = add_lag_features(df)
    df = add_difference_features(df)
    df['SMA_7_Squared'] = df['7-Day Moving Average'] ** 2
    df['SMA_30_Squared'] = df['30-Day Moving Average'] ** 2
    df['RSI_Squared'] = df['RSI_14'] ** 2
    df['MACD_Squared'] = df['MACD'] ** 2
    df['Volatility_Squared'] = df['Rolling Volatility (30d)'] ** 2
    df['RSI_x_Volume'] = df['RSI_14'] * df['Volume']
    df['MACD_x_Price'] = df['MACD'] * df['Close']
    df['SMA30_x_Volatility'] = df['30-Day Moving Average'] * df['Rolling_
↳Volatility (30d)']
    df['VolumePriceRatio_x_Return'] = df['Volume_Price_Ratio'] * df['Daily_
↳Return']
    return df

filtered_df_sorted = full_feature_engineering(filtered_df_sorted)
filtered_df_sorted['level_0'] = filtered_df_sorted['Ticker']

```

3.2 MODELLING

I applied SARIMAX (by incorporating seasonality & exogenous inputs to predict AMZN closing prices), Ensemble Model (by taking average or weighted predictions from Week 2 models (Linear Regression, ARIMA, Random Forest) using their MAE scores) and Temporal Fusion Transformer (exploring the TFT through provided readings).

3.2.1 SARIMAX

```
[130]: from scipy.stats import boxcox, jarque_bera
from arch import arch_model
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from pmdarima.model_selection import SlidingWindowForecastCV
from pmdarima import auto_arima
from statsmodels.stats.diagnostic import acorr_ljungbox
from sklearn.metrics import accuracy_score
import warnings
warnings.filterwarnings("ignore")
def inv_boxcox(y, lambda_val):
    with np.errstate(over='ignore', invalid='ignore'):
        if lambda_val == 0:
            result = np.exp(y)
        else:
            result = (lambda_val * y + 1) ** (1 / lambda_val)
        result = np.nan_to_num(result, nan=np.nanmedian(result),
                                posinf=np.nanmax(result), neginf=np.
    nanmin(result))
    return result
def safe_predict(model, n_periods, X, lambda_val):
    try:
        forecast, _ = model.predict(
            n_periods=n_periods,
            X=X,
            return_conf_int=True
        )
    except Exception as e:
        print(f"Prediction error: {str(e)}")
        forecast = np.full(n_periods, np.nan)
        forecast = pd.Series(forecast).replace([np.inf, -np.inf], np.nan)
        forecast = inv_boxcox(forecast.values, lambda_val)
    return forecast
if 'level_0' in filtered_df_sorted.columns:
    filtered_df_sorted.drop(columns=['level_0'], inplace=True)
if not isinstance(filtered_df_sorted.index, pd.MultiIndex):
    filtered_df_sorted = filtered_df_sorted.set_index(['Ticker', 'Date'])
amzn_df = filtered_df_sorted.loc['amzn'].copy().reset_index()
amzn_df = amzn_df[amzn_df['Date'] >= '2010-01-04'].set_index('Date')
amzn_df['SP500_Return'] = amzn_df['SP500_Close'].pct_change().clip(-0.1, 0.1)
```

```

amzn_df['SP500_Lag1'] = amzn_df['SP500_Return'].shift(1)
amzn_df['VIX_Lag1'] = amzn_df['VIX'].shift(1)
amzn_df['Bollinger_Width'] = (
    (amzn_df['30-Day Moving Average'] + 2*amzn_df['Rolling Volatility (30d)']) -
    (amzn_df['30-Day Moving Average'] - 2*amzn_df['Rolling Volatility (30d)'])
)
for col in ['Close_trend', 'Close_seasonal', 'RSI_x_Volume',
            'RSI_14', 'MACD', 'Bollinger_Width']:
    amzn_df[col] = amzn_df[col].ffill().bfill().fillna(0)
train_size = int(len(amzn_df) * 0.7)
val_size = int(len(amzn_df) * 0.15)
test_size = len(amzn_df) - train_size - val_size
train_df = amzn_df.iloc[:train_size]
val_df = amzn_df.iloc[train_size:train_size+val_size]
test_df = amzn_df.iloc[train_size+val_size:]
exog_vars = [
    'SP500_Return', 'SP500_Lag1', 'VIX', 'VIX_Lag1',
    'Close_trend', 'Close_seasonal', 'RSI_x_Volume',
    'RSI_14', 'MACD', 'Bollinger_Width'
]
def train_sarimax(train_data, exog_data):
    train_data = train_data.copy()
    exog_data = exog_data[exog_vars].ffill().bfill().fillna(0)
    train_data['Close_Trans'], lambda_val = boxcox(
        train_data['Close'].clip(lower=1e-6)
    )
    model = auto_arima(
        y=train_data['Close_Trans'],
        X=exog_data,
        seasonal=True,
        m=12,
        d=1,
        D=1,
        start_p=0, max_p=4,
        start_q=0, max_q=4,
        start_P=0, max_P=3,
        start_Q=0, max_Q=2,
        information_criterion='bic',
        trace=True,
        error_action='ignore',
        suppress_warnings=True,
        stepwise=True,
        maxiter=50
    )
    return model, lambda_val
def run_cross_validation():
    cv = SlidingWindowForecastCV(window_size=252, step=63, h=21)

```

```

cv_gen = cv.split(amzn_df)
scores = []
fold_count = 0
total_folds = sum(1 for _ in cv.split(amzn_df))
cv_gen = cv.split(amzn_df)
for train_idx, test_idx in cv_gen:
    fold_count += 1
    if len(train_idx) < 200 or len(test_idx) < 10:
        print(f"Skipping fold {fold_count}/{total_folds} - insufficient_
↳data")
        continue
    print(f"\nCV Fold {fold_count}/{total_folds}")
    train_data = amzn_df.iloc[train_idx]
    test_data = amzn_df.iloc[test_idx]
    try:
        model, lambda_val = train_sarimax(train_data, train_data)
        forecast = safe_predict(model, len(test_data),
↳test_data[exog_vars], lambda_val)
        actual = test_data['Close'].replace([np.inf, -np.inf], np.nan)
        if actual.isna().any():
            actual = actual.ffill().bfill()
        if len(forecast) > 1 and len(actual) > 1:
            actual_dir = actual.diff().dropna() > 0
            pred_dir = pd.Series(forecast).diff().dropna() > 0
            common_idx = actual_dir.index.intersection(pred_dir.index)
            if len(common_idx) > 10:
                accuracy = accuracy_score(
                    actual_dir.loc[common_idx],
                    pred_dir.loc[common_idx]
                )
                scores.append(accuracy)
                print(f"Fold Accuracy: {accuracy:.4f}")
    except Exception as e:
        print(f"CV fold failed: {str(e)}")
    return np.mean(scores) if scores else 0
cv_score = run_cross_validation()
print(f"\nFinal Cross-Validation Accuracy: {cv_score:.4f}")
combined_df = pd.concat([train_df, val_df])
final_model, lambda_val = train_sarimax(combined_df, combined_df)
residuals = final_model.arima_res_.resid.dropna()
if len(residuals) > 10:
    garch_model = arch_model(residuals, vol='Garch', p=1, q=1)
    garch_fit = garch_model.fit(disp='off')
    print(f"GARCH model fitted. AIC: {garch_fit.aic:.2f}")
else:
    print("Insufficient residuals for GARCH modeling")
    garch_fit = None

```

```

test_exog = test_df[exog_vars].ffill().bfill().fillna(0)
test_df['Forecast_Close'] = safe_predict(
    final_model,
    len(test_df),
    test_exog,
    lambda_val
)
if test_df['Forecast_Close'].isna().any():
    print(f"Replacing {test_df['Forecast_Close'].isna().sum()} NaN forecasts")
    last_train = combined_df['Close'].iloc[-1]
    test_df['Forecast_Close'] = test_df['Forecast_Close'].fillna(last_train)
test_df['Actual_Close'] = test_df['Close'].replace([np.inf, -np.inf], np.nan)
if test_df['Actual_Close'].isna().any():
    test_df['Actual_Close'] = test_df['Actual_Close'].ffill().bfill()
valid_mask = test_df[['Actual_Close', 'Forecast_Close']].notnull().all(axis=1)
test_df = test_df[valid_mask]
if len(test_df) > 0:
    mae_sarimax = mean_absolute_error(test_df['Actual_Close'],
    ↪test_df['Forecast_Close'])
    print(f"MAE: {mae_sarimax:.4f}")
    test_df['Actual_Direction'] = test_df['Actual_Close'].diff().gt(0)
    test_df['Predicted_Direction'] = test_df['Forecast_Close'].diff().gt(0)
    test_df.dropna(subset=['Actual_Direction', 'Predicted_Direction'],
    ↪inplace=True)
    if len(test_df) > 0:
        accuracy_sarimax = accuracy_score(
            test_df['Actual_Direction'],
            test_df['Predicted_Direction']
        )
        print(f"Directional Accuracy: {accuracy_sarimax:.4f}")
    else:
        print("Insufficient data for directional accuracy")
else:
    print("No valid data for evaluation")
def run_diagnostics(model):
    try:
        residuals = model.arma_res_.resid.dropna()
        if len(residuals) == 0:
            print("No residuals for diagnostics")
            return
        fig, axes = plt.subplots(3, 1, figsize=(12, 12))
        axes[0].plot(residuals)
        axes[0].set_title('Residuals over Time')
        axes[0].grid(alpha=0.3)
        sns.histplot(residuals, kde=True, ax=axes[1])
        axes[1].set_title('Residual Distribution')
        plot_acf(residuals, lags=40, ax=axes[2])
    
```



```

axes[2].set_title('Residual ACF')
plt.tight_layout()
plt.show()
if len(residuals) > 10:
    lb_test = acorr_ljungbox(residuals, lags=[10], return_df=True)
    jb_test = jarque_bera(residuals)
    print(f"Ljung-Box p-value: {lb_test['lb_pvalue'].iloc[0]:.4f}")
    print(f"Jarque-Bera p-value: {jb_test[1]:.4f}")
except Exception as e:
    print(f"Diagnostics failed: {str(e)}")
print("\nRunning model diagnostics...")
run_diagnostics(final_model)

```

CV Fold 1/21

Performing stepwise search to minimize bic

ARIMA(0,1,0)(0,1,0)[12]	: BIC=-4793.607, Time=0.40 sec
ARIMA(1,1,0)(1,1,0)[12]	: BIC=-592.388, Time=1.84 sec
ARIMA(0,1,1)(0,1,1)[12]	: BIC=-263.698, Time=2.12 sec
ARIMA(0,1,0)(1,1,0)[12]	: BIC=-332.998, Time=3.80 sec
ARIMA(0,1,0)(0,1,1)[12]	: BIC=-162.600, Time=1.69 sec
ARIMA(0,1,0)(1,1,1)[12]	: BIC=-466.219, Time=1.85 sec
ARIMA(1,1,0)(0,1,0)[12]	: BIC=-4788.098, Time=0.48 sec
ARIMA(0,1,1)(0,1,0)[12]	: BIC=-4788.868, Time=0.79 sec
ARIMA(1,1,1)(0,1,0)[12]	: BIC=-4832.920, Time=0.87 sec
ARIMA(1,1,1)(1,1,0)[12]	: BIC=-528.610, Time=1.98 sec
ARIMA(1,1,1)(0,1,1)[12]	: BIC=-255.065, Time=3.42 sec
ARIMA(1,1,1)(1,1,1)[12]	: BIC=-134.498, Time=3.61 sec
ARIMA(2,1,1)(0,1,0)[12]	: BIC=-3619.999, Time=0.98 sec
ARIMA(1,1,2)(0,1,0)[12]	: BIC=-3993.296, Time=0.90 sec
ARIMA(0,1,2)(0,1,0)[12]	: BIC=-2920.826, Time=0.88 sec
ARIMA(2,1,0)(0,1,0)[12]	: BIC=-412.506, Time=1.04 sec
ARIMA(2,1,2)(0,1,0)[12]	: BIC=-3963.775, Time=0.95 sec
ARIMA(1,1,1)(0,1,0)[12] intercept	: BIC=-4827.449, Time=0.91 sec

Best model: ARIMA(1,1,1)(0,1,0)[12]

Total fit time: 28.523 seconds

CV Fold 2/21

Performing stepwise search to minimize bic

ARIMA(0,1,0)(0,1,0)[12]	: BIC=2413.495, Time=0.44 sec
ARIMA(1,1,0)(1,1,0)[12]	: BIC=1308.423, Time=1.63 sec
ARIMA(0,1,1)(0,1,1)[12]	: BIC=1281.381, Time=4.02 sec
ARIMA(0,1,1)(0,1,0)[12]	: BIC=1348.053, Time=0.79 sec
ARIMA(0,1,1)(1,1,1)[12]	: BIC=1298.454, Time=1.98 sec
ARIMA(0,1,1)(0,1,2)[12]	: BIC=1279.673, Time=3.87 sec
ARIMA(0,1,1)(1,1,2)[12]	: BIC=1285.641, Time=6.16 sec

ARIMA(0,1,0)(0,1,2)[12]	: BIC=1257.295, Time=3.25 sec
ARIMA(0,1,0)(0,1,1)[12]	: BIC=1264.946, Time=1.51 sec
ARIMA(0,1,0)(1,1,2)[12]	: BIC=1258.073, Time=3.47 sec
ARIMA(0,1,0)(1,1,1)[12]	: BIC=1282.753, Time=1.66 sec
ARIMA(1,1,0)(0,1,2)[12]	: BIC=1278.706, Time=5.61 sec
ARIMA(1,1,1)(0,1,2)[12]	: BIC=1285.136, Time=4.23 sec
ARIMA(0,1,0)(0,1,2)[12] intercept	: BIC=1279.636, Time=3.52 sec

Best model: ARIMA(0,1,0)(0,1,2)[12]

Total fit time: 42.156 seconds

CV Fold 3/21

Performing stepwise search to minimize bic

ARIMA(0,1,0)(0,1,0)[12]	: BIC=8073.026, Time=0.38 sec
ARIMA(1,1,0)(1,1,0)[12]	: BIC=6957.547, Time=3.28 sec
ARIMA(0,1,1)(0,1,1)[12]	: BIC=6917.138, Time=3.76 sec
ARIMA(0,1,1)(0,1,0)[12]	: BIC=7007.496, Time=0.37 sec
ARIMA(0,1,1)(1,1,1)[12]	: BIC=6922.475, Time=5.47 sec
ARIMA(0,1,1)(0,1,2)[12]	: BIC=6922.461, Time=10.56 sec
ARIMA(0,1,1)(1,1,0)[12]	: BIC=6957.252, Time=1.16 sec
ARIMA(0,1,1)(1,1,2)[12]	: BIC=6928.964, Time=8.61 sec
ARIMA(0,1,0)(0,1,1)[12]	: BIC=6893.016, Time=2.39 sec
ARIMA(0,1,0)(1,1,1)[12]	: BIC=6895.488, Time=6.22 sec
ARIMA(0,1,0)(0,1,2)[12]	: BIC=6914.514, Time=5.29 sec
ARIMA(0,1,0)(1,1,0)[12]	: BIC=6944.685, Time=2.12 sec
ARIMA(0,1,0)(1,1,2)[12]	: BIC=6905.431, Time=10.04 sec
ARIMA(1,1,0)(0,1,1)[12]	: BIC=6917.618, Time=2.20 sec
ARIMA(1,1,1)(0,1,1)[12]	: BIC=6925.030, Time=2.12 sec
ARIMA(0,1,0)(0,1,1)[12] intercept	: BIC=6919.551, Time=2.15 sec

Best model: ARIMA(0,1,0)(0,1,1)[12]

Total fit time: 66.150 seconds

CV Fold 4/21

Performing stepwise search to minimize bic

ARIMA(0,1,0)(0,1,0)[12]	: BIC=-3428.905, Time=0.40 sec
ARIMA(1,1,0)(1,1,0)[12]	: BIC=-4546.716, Time=3.72 sec
ARIMA(0,1,1)(0,1,1)[12]	: BIC=-4576.309, Time=1.81 sec
ARIMA(0,1,1)(0,1,0)[12]	: BIC=-4494.644, Time=0.77 sec
ARIMA(0,1,1)(1,1,1)[12]	: BIC=-4566.541, Time=1.98 sec
ARIMA(0,1,1)(0,1,2)[12]	: BIC=-4571.245, Time=3.98 sec
ARIMA(0,1,1)(1,1,0)[12]	: BIC=46786869.848, Time=1.98 sec
ARIMA(0,1,1)(1,1,2)[12]	: BIC=-4565.136, Time=5.84 sec
ARIMA(0,1,0)(0,1,1)[12]	: BIC=-3767.955, Time=1.53 sec
ARIMA(1,1,1)(0,1,1)[12]	: BIC=-4570.901, Time=1.94 sec
ARIMA(0,1,2)(0,1,1)[12]	: BIC=-4570.553, Time=1.84 sec
ARIMA(1,1,0)(0,1,1)[12]	: BIC=-4575.403, Time=1.62 sec
ARIMA(1,1,2)(0,1,1)[12]	: BIC=-4566.193, Time=4.21 sec

ARIMA(0,1,1)(0,1,1)[12] intercept : BIC=-4571.368, Time=1.97 sec

Best model: ARIMA(0,1,1)(0,1,1)[12]

Total fit time: 33.612 seconds

CV Fold 5/21

Performing stepwise search to minimize bic

ARIMA(0,1,0)(0,1,0)[12]	: BIC=-3038.248, Time=0.40 sec
ARIMA(1,1,0)(1,1,0)[12]	: BIC=-4176.441, Time=1.64 sec
ARIMA(0,1,1)(0,1,1)[12]	: BIC=-4203.322, Time=1.81 sec
ARIMA(0,1,1)(0,1,0)[12]	: BIC=-4103.660, Time=0.76 sec
ARIMA(0,1,1)(1,1,1)[12]	: BIC=-4192.527, Time=2.00 sec
ARIMA(0,1,1)(0,1,2)[12]	: BIC=-4198.787, Time=6.03 sec
ARIMA(0,1,1)(1,1,0)[12]	: BIC=-4175.324, Time=1.57 sec
ARIMA(0,1,1)(1,1,2)[12]	: BIC=-4193.446, Time=4.26 sec
ARIMA(0,1,0)(0,1,1)[12]	: BIC=-4244.099, Time=1.48 sec
ARIMA(0,1,0)(1,1,1)[12]	: BIC=-4232.621, Time=3.83 sec
ARIMA(0,1,0)(0,1,2)[12]	: BIC=-4235.553, Time=3.24 sec
ARIMA(0,1,0)(1,1,0)[12]	: BIC=-4201.481, Time=1.44 sec
ARIMA(0,1,0)(1,1,2)[12]	: BIC=-4244.603, Time=3.51 sec
ARIMA(0,1,0)(2,1,2)[12]	: BIC=-4237.208, Time=5.86 sec
ARIMA(0,1,0)(2,1,1)[12]	: BIC=-4224.916, Time=3.50 sec
ARIMA(1,1,0)(1,1,2)[12]	: BIC=-4194.853, Time=3.79 sec
ARIMA(1,1,1)(1,1,2)[12]	: BIC=-4188.580, Time=6.59 sec
ARIMA(0,1,0)(1,1,2)[12] intercept	: BIC=-4195.544, Time=3.79 sec

Best model: ARIMA(0,1,0)(1,1,2)[12]

Total fit time: 55.516 seconds

CV Fold 6/21

Performing stepwise search to minimize bic

ARIMA(0,1,0)(0,1,0)[12]	: BIC=-4793.612, Time=0.39 sec
ARIMA(1,1,0)(1,1,0)[12]	: BIC=-570.860, Time=1.85 sec
ARIMA(0,1,1)(0,1,1)[12]	: BIC=-4542.273, Time=1.86 sec
ARIMA(0,1,0)(1,1,0)[12]	: BIC=-358.739, Time=3.70 sec
ARIMA(0,1,0)(0,1,1)[12]	: BIC=-4545.281, Time=1.63 sec
ARIMA(0,1,0)(1,1,1)[12]	: BIC=-4550.411, Time=1.66 sec
ARIMA(1,1,0)(0,1,0)[12]	: BIC=-1818.229, Time=0.54 sec
ARIMA(0,1,1)(0,1,0)[12]	: BIC=-4902.786, Time=0.78 sec
ARIMA(0,1,1)(1,1,0)[12]	: BIC=-567.001, Time=1.78 sec
ARIMA(0,1,1)(1,1,1)[12]	: BIC=-1948.430, Time=2.15 sec
ARIMA(1,1,1)(0,1,0)[12]	: BIC=-4896.660, Time=0.86 sec
ARIMA(0,1,2)(0,1,0)[12]	: BIC=-4873.635, Time=0.87 sec
ARIMA(1,1,2)(0,1,0)[12]	: BIC=-3424.071, Time=2.99 sec
ARIMA(0,1,1)(0,1,0)[12] intercept	: BIC=-4897.308, Time=1.08 sec

Best model: ARIMA(0,1,1)(0,1,0)[12]

Total fit time: 22.168 seconds

CV Fold 7/21

Performing stepwise search to minimize bic

ARIMA(0,1,0)(0,1,0)[12]	: BIC=117.563, Time=0.39 sec
ARIMA(1,1,0)(1,1,0)[12]	: BIC=-1022.002, Time=1.65 sec
ARIMA(0,1,1)(0,1,1)[12]	: BIC=-1048.942, Time=1.82 sec
ARIMA(0,1,1)(0,1,0)[12]	: BIC=-947.862, Time=0.80 sec
ARIMA(0,1,1)(1,1,1)[12]	: BIC=-1038.545, Time=2.01 sec
ARIMA(0,1,1)(0,1,2)[12]	: BIC=-1045.777, Time=5.68 sec
ARIMA(0,1,1)(1,1,0)[12]	: BIC=-1021.322, Time=1.69 sec
ARIMA(0,1,1)(1,1,2)[12]	: BIC=-1040.647, Time=4.27 sec
ARIMA(0,1,0)(0,1,1)[12]	: BIC=-1094.152, Time=1.49 sec
ARIMA(0,1,0)(1,1,1)[12]	: BIC=-1082.226, Time=1.64 sec
ARIMA(0,1,0)(0,1,2)[12]	: BIC=-1096.199, Time=5.19 sec
ARIMA(0,1,0)(1,1,2)[12]	: BIC=-1098.461, Time=3.50 sec
ARIMA(0,1,0)(2,1,2)[12]	: BIC=-1075.273, Time=3.87 sec
ARIMA(0,1,0)(2,1,1)[12]	: BIC=-1072.295, Time=5.57 sec
ARIMA(1,1,0)(1,1,2)[12]	: BIC=-1041.504, Time=3.75 sec
ARIMA(1,1,1)(1,1,2)[12]	: BIC=-1035.441, Time=4.50 sec
ARIMA(0,1,0)(1,1,2)[12] intercept	: BIC=-1041.451, Time=5.89 sec

Best model: ARIMA(0,1,0)(1,1,2)[12]

Total fit time: 53.735 seconds

CV Fold 8/21

Performing stepwise search to minimize bic

ARIMA(0,1,0)(0,1,0)[12]	: BIC=-1429.741, Time=0.38 sec
ARIMA(1,1,0)(1,1,0)[12]	: BIC=-2557.990, Time=1.63 sec
ARIMA(0,1,1)(0,1,1)[12]	: BIC=-2584.730, Time=1.81 sec
ARIMA(0,1,1)(0,1,0)[12]	: BIC=-2495.573, Time=0.79 sec
ARIMA(0,1,1)(1,1,1)[12]	: BIC=-2573.993, Time=2.01 sec
ARIMA(0,1,1)(0,1,2)[12]	: BIC=-2579.961, Time=5.79 sec
ARIMA(0,1,1)(1,1,0)[12]	: BIC=-2558.574, Time=1.57 sec
ARIMA(0,1,1)(1,1,2)[12]	: BIC=-2574.032, Time=4.27 sec
ARIMA(0,1,0)(0,1,1)[12]	: BIC=-2617.549, Time=1.48 sec
ARIMA(0,1,0)(1,1,1)[12]	: BIC=-2596.582, Time=1.66 sec
ARIMA(0,1,0)(0,1,2)[12]	: BIC=-2605.761, Time=5.37 sec
ARIMA(0,1,0)(1,1,0)[12]	: BIC=-2575.428, Time=1.48 sec
ARIMA(0,1,0)(1,1,2)[12]	: BIC=-2582.311, Time=3.53 sec
ARIMA(1,1,0)(0,1,1)[12]	: BIC=-2584.155, Time=1.58 sec
ARIMA(1,1,1)(0,1,1)[12]	: BIC=-2579.257, Time=2.55 sec
ARIMA(0,1,0)(0,1,1)[12] intercept	: BIC=-2584.861, Time=3.24 sec

Best model: ARIMA(0,1,0)(0,1,1)[12]

Total fit time: 39.152 seconds

CV Fold 9/21

Performing stepwise search to minimize bic

ARIMA(0,1,0)(0,1,0)[12]	: BIC=4571.798, Time=0.39 sec
ARIMA(1,1,0)(1,1,0)[12]	: BIC=3455.884, Time=3.99 sec
ARIMA(0,1,1)(0,1,1)[12]	: BIC=3433.416, Time=4.15 sec
ARIMA(0,1,1)(0,1,0)[12]	: BIC=3506.167, Time=4.23 sec
ARIMA(0,1,1)(1,1,1)[12]	: BIC=3445.449, Time=4.76 sec
ARIMA(0,1,1)(0,1,2)[12]	: BIC=3438.027, Time=4.36 sec
ARIMA(0,1,1)(1,1,0)[12]	: BIC=3456.025, Time=6.01 sec
ARIMA(0,1,1)(1,1,2)[12]	: BIC=3443.827, Time=4.33 sec
ARIMA(0,1,0)(0,1,1)[12]	: BIC=3404.060, Time=1.48 sec
ARIMA(0,1,0)(1,1,1)[12]	: BIC=3419.956, Time=1.63 sec
ARIMA(0,1,0)(0,1,2)[12]	: BIC=3405.319, Time=5.35 sec
ARIMA(0,1,0)(1,1,0)[12]	: BIC=3439.963, Time=3.49 sec
ARIMA(0,1,0)(1,1,2)[12]	: BIC=3409.949, Time=3.48 sec
ARIMA(1,1,0)(0,1,1)[12]	: BIC=3433.144, Time=5.92 sec
ARIMA(1,1,1)(0,1,1)[12]	: BIC=3438.998, Time=1.95 sec
ARIMA(0,1,0)(0,1,1)[12] intercept	: BIC=3433.866, Time=1.60 sec

Best model: ARIMA(0,1,0)(0,1,1)[12]

Total fit time: 57.137 seconds

CV Fold 10/21

Performing stepwise search to minimize bic

ARIMA(0,1,0)(0,1,0)[12]	: BIC=8195.467, Time=0.25 sec
ARIMA(1,1,0)(1,1,0)[12]	: BIC=7130.527, Time=2.69 sec
ARIMA(0,1,1)(0,1,1)[12]	: BIC=7130.011, Time=3.77 sec
ARIMA(0,1,1)(0,1,0)[12]	: BIC=7129.789, Time=2.14 sec
ARIMA(0,1,1)(1,1,0)[12]	: BIC=7130.553, Time=2.78 sec
ARIMA(0,1,1)(1,1,1)[12]	: BIC=inf, Time=3.97 sec
ARIMA(1,1,1)(0,1,0)[12]	: BIC=inf, Time=2.56 sec
ARIMA(0,1,2)(0,1,0)[12]	: BIC=7135.251, Time=0.79 sec
ARIMA(1,1,0)(0,1,0)[12]	: BIC=7129.786, Time=1.66 sec
ARIMA(1,1,0)(0,1,1)[12]	: BIC=7129.983, Time=3.38 sec
ARIMA(1,1,0)(1,1,1)[12]	: BIC=inf, Time=3.93 sec
ARIMA(2,1,0)(0,1,0)[12]	: BIC=7135.257, Time=0.97 sec
ARIMA(2,1,1)(0,1,0)[12]	: BIC=7140.495, Time=1.16 sec
ARIMA(1,1,0)(0,1,0)[12] intercept	: BIC=7135.216, Time=0.92 sec

Best model: ARIMA(1,1,0)(0,1,0)[12]

Total fit time: 30.984 seconds

CV Fold 11/21

Performing stepwise search to minimize bic

ARIMA(0,1,0)(0,1,0)[12]	: BIC=12105.532, Time=0.05 sec
ARIMA(1,1,0)(1,1,0)[12]	: BIC=11036.226, Time=0.73 sec
ARIMA(0,1,1)(0,1,1)[12]	: BIC=11032.711, Time=2.90 sec
ARIMA(0,1,1)(0,1,0)[12]	: BIC=11040.099, Time=0.79 sec
ARIMA(0,1,1)(1,1,1)[12]	: BIC=inf, Time=2.25 sec
ARIMA(0,1,1)(0,1,2)[12]	: BIC=11015.106, Time=5.09 sec

ARIMA(0,1,1)(1,1,2)[12]	: BIC=11014.364, Time=7.69 sec
ARIMA(0,1,1)(2,1,2)[12]	: BIC=11019.651, Time=11.84 sec
ARIMA(0,1,1)(2,1,1)[12]	: BIC=11014.293, Time=3.88 sec
ARIMA(0,1,1)(2,1,0)[12]	: BIC=11027.687, Time=8.11 sec
ARIMA(0,1,1)(3,1,1)[12]	: BIC=11019.681, Time=11.08 sec
ARIMA(0,1,1)(1,1,0)[12]	: BIC=11036.304, Time=0.64 sec
ARIMA(0,1,1)(3,1,0)[12]	: BIC=11025.721, Time=5.13 sec
ARIMA(0,1,1)(3,1,2)[12]	: BIC=11025.237, Time=6.93 sec
ARIMA(0,1,0)(2,1,1)[12]	: BIC=inf, Time=5.98 sec
ARIMA(1,1,1)(2,1,1)[12]	: BIC=11019.208, Time=4.08 sec
ARIMA(0,1,2)(2,1,1)[12]	: BIC=11018.470, Time=4.22 sec
ARIMA(1,1,0)(2,1,1)[12]	: BIC=11013.934, Time=7.83 sec
ARIMA(1,1,0)(1,1,1)[12]	: BIC=inf, Time=1.98 sec
ARIMA(1,1,0)(2,1,0)[12]	: BIC=11027.501, Time=3.51 sec
ARIMA(1,1,0)(3,1,1)[12]	: BIC=11019.326, Time=10.75 sec
ARIMA(1,1,0)(2,1,2)[12]	: BIC=11019.296, Time=11.44 sec
ARIMA(1,1,0)(1,1,2)[12]	: BIC=11014.010, Time=4.03 sec
ARIMA(1,1,0)(3,1,0)[12]	: BIC=11025.476, Time=5.54 sec
ARIMA(1,1,0)(3,1,2)[12]	: BIC=11024.871, Time=10.04 sec
ARIMA(2,1,0)(2,1,1)[12]	: BIC=11017.945, Time=6.82 sec
ARIMA(2,1,1)(2,1,1)[12]	: BIC=11023.006, Time=10.44 sec
ARIMA(1,1,0)(2,1,1)[12] intercept	: BIC=11019.109, Time=5.85 sec

Best model: ARIMA(1,1,0)(2,1,1)[12]

Total fit time: 159.667 seconds

CV Fold 12/21

Performing stepwise search to minimize bic

ARIMA(0,1,0)(0,1,0)[12]	: BIC=-894.326, Time=0.37 sec
ARIMA(1,1,0)(1,1,0)[12]	: BIC=-2012.932, Time=1.63 sec
ARIMA(0,1,1)(0,1,1)[12]	: BIC=-2046.446, Time=1.83 sec
ARIMA(0,1,1)(0,1,0)[12]	: BIC=-1960.072, Time=2.54 sec
ARIMA(0,1,1)(1,1,1)[12]	: BIC=-2037.950, Time=2.55 sec
ARIMA(0,1,1)(0,1,2)[12]	: BIC=inf, Time=3.96 sec
ARIMA(0,1,1)(1,1,0)[12]	: BIC=-2013.116, Time=1.58 sec
ARIMA(0,1,1)(1,1,2)[12]	: BIC=inf, Time=6.24 sec
ARIMA(0,1,0)(0,1,1)[12]	: BIC=-2080.130, Time=1.52 sec
ARIMA(0,1,0)(1,1,1)[12]	: BIC=-2068.530, Time=1.63 sec
ARIMA(0,1,0)(0,1,2)[12]	: BIC=inf, Time=3.23 sec
ARIMA(0,1,0)(1,1,0)[12]	: BIC=-2027.191, Time=1.46 sec
ARIMA(0,1,0)(1,1,2)[12]	: BIC=inf, Time=5.48 sec
ARIMA(1,1,0)(0,1,1)[12]	: BIC=-2045.922, Time=1.62 sec
ARIMA(1,1,1)(0,1,1)[12]	: BIC=-2044.056, Time=1.92 sec
ARIMA(0,1,0)(0,1,1)[12] intercept	: BIC=-2045.097, Time=1.62 sec

Best model: ARIMA(0,1,0)(0,1,1)[12]

Total fit time: 39.214 seconds

CV Fold 13/21

Performing stepwise search to minimize bic

ARIMA(0,1,0)(0,1,0)[12]	: BIC=-4793.612, Time=0.40 sec
ARIMA(1,1,0)(1,1,0)[12]	: BIC=-167.867, Time=1.82 sec
ARIMA(0,1,1)(0,1,1)[12]	: BIC=-157.443, Time=2.07 sec
ARIMA(0,1,0)(1,1,0)[12]	: BIC=-230.611, Time=3.79 sec
ARIMA(0,1,0)(0,1,1)[12]	: BIC=-201.789, Time=1.67 sec
ARIMA(0,1,0)(1,1,1)[12]	: BIC=-298.764, Time=1.85 sec
ARIMA(1,1,0)(0,1,0)[12]	: BIC=-4914.178, Time=0.47 sec
ARIMA(1,1,0)(0,1,1)[12]	: BIC=-196.572, Time=1.79 sec
ARIMA(1,1,0)(1,1,1)[12]	: BIC=-511.679, Time=2.01 sec
ARIMA(2,1,0)(0,1,0)[12]	: BIC=-3700.052, Time=0.91 sec
ARIMA(1,1,1)(0,1,0)[12]	: BIC=-3817.098, Time=0.91 sec
ARIMA(0,1,1)(0,1,0)[12]	: BIC=-4789.168, Time=2.53 sec
ARIMA(2,1,1)(0,1,0)[12]	: BIC=-350.473, Time=1.69 sec
ARIMA(1,1,0)(0,1,0)[12] intercept	: BIC=-4908.702, Time=0.49 sec

Best model: ARIMA(1,1,0)(0,1,0)[12]

Total fit time: 22.437 seconds

CV Fold 14/21

Performing stepwise search to minimize bic

ARIMA(0,1,0)(0,1,0)[12]	: BIC=-642.528, Time=0.39 sec
ARIMA(1,1,0)(1,1,0)[12]	: BIC=-1777.162, Time=1.62 sec
ARIMA(0,1,1)(0,1,1)[12]	: BIC=-1809.409, Time=1.82 sec
ARIMA(0,1,1)(0,1,0)[12]	: BIC=-1707.856, Time=0.80 sec
ARIMA(0,1,1)(1,1,1)[12]	: BIC=-1798.171, Time=1.99 sec
ARIMA(0,1,1)(0,1,2)[12]	: BIC=-1804.987, Time=6.03 sec
ARIMA(0,1,1)(1,1,0)[12]	: BIC=-1779.494, Time=1.56 sec
ARIMA(0,1,1)(1,1,2)[12]	: BIC=-1799.243, Time=4.29 sec
ARIMA(0,1,0)(0,1,1)[12]	: BIC=-1848.288, Time=1.47 sec
ARIMA(0,1,0)(1,1,1)[12]	: BIC=-1834.489, Time=1.65 sec
ARIMA(0,1,0)(0,1,2)[12]	: BIC=-1839.697, Time=5.35 sec
ARIMA(0,1,0)(1,1,0)[12]	: BIC=-1800.687, Time=1.44 sec
ARIMA(0,1,0)(1,1,2)[12]	: BIC=-1843.555, Time=3.51 sec
ARIMA(1,1,0)(0,1,1)[12]	: BIC=-1805.014, Time=1.60 sec
ARIMA(1,1,1)(0,1,1)[12]	: BIC=-1805.600, Time=2.48 sec
ARIMA(0,1,0)(0,1,1)[12] intercept	: BIC=-1805.635, Time=3.31 sec

Best model: ARIMA(0,1,0)(0,1,1)[12]

Total fit time: 39.316 seconds

CV Fold 15/21

Performing stepwise search to minimize bic

ARIMA(0,1,0)(0,1,0)[12]	: BIC=-4464.149, Time=0.40 sec
ARIMA(1,1,0)(1,1,0)[12]	: BIC=-388.062, Time=1.91 sec
ARIMA(0,1,1)(0,1,1)[12]	: BIC=inf, Time=1.85 sec
ARIMA(0,1,0)(1,1,0)[12]	: BIC=-647.263, Time=1.61 sec

ARIMA(0,1,0)(0,1,1)[12]	: BIC=inf, Time=1.54 sec
ARIMA(0,1,0)(1,1,1)[12]	: BIC=-890.542, Time=1.87 sec
ARIMA(1,1,0)(0,1,0)[12]	: BIC=-280.156, Time=0.93 sec
ARIMA(0,1,1)(0,1,0)[12]	: BIC=8465.081, Time=2.39 sec
ARIMA(1,1,1)(0,1,0)[12]	: BIC=586.765, Time=1.04 sec
ARIMA(0,1,0)(0,1,0)[12] intercept	: BIC=-4778.808, Time=0.46 sec
ARIMA(0,1,0)(1,1,0)[12] intercept	: BIC=-641.679, Time=1.78 sec
ARIMA(0,1,0)(0,1,1)[12] intercept	: BIC=inf, Time=1.64 sec
ARIMA(0,1,0)(1,1,1)[12] intercept	: BIC=-897.828, Time=1.96 sec
ARIMA(1,1,0)(0,1,0)[12] intercept	: BIC=-280.224, Time=0.62 sec
ARIMA(0,1,1)(0,1,0)[12] intercept	: BIC=5341.123, Time=0.97 sec
ARIMA(1,1,1)(0,1,0)[12] intercept	: BIC=13728.473, Time=1.09 sec

Best model: ARIMA(0,1,0)(0,1,0)[12] intercept
Total fit time: 22.047 seconds

CV Fold 16/21

Performing stepwise search to minimize bic

ARIMA(0,1,0)(0,1,0)[12]	: BIC=-3984.355, Time=0.64 sec
ARIMA(1,1,0)(1,1,0)[12]	: BIC=-312.781, Time=3.56 sec
ARIMA(0,1,1)(0,1,1)[12]	: BIC=-509.449, Time=2.02 sec
ARIMA(0,1,0)(1,1,0)[12]	: BIC=-302.340, Time=1.60 sec
ARIMA(0,1,0)(0,1,1)[12]	: BIC=-359.815, Time=1.73 sec
ARIMA(0,1,0)(1,1,1)[12]	: BIC=-490.535, Time=1.86 sec
ARIMA(1,1,0)(0,1,0)[12]	: BIC=-4815.283, Time=0.46 sec
ARIMA(1,1,0)(0,1,1)[12]	: BIC=-508.542, Time=3.19 sec
ARIMA(1,1,0)(1,1,1)[12]	: BIC=-711.492, Time=2.95 sec
ARIMA(2,1,0)(0,1,0)[12]	: BIC=-781.578, Time=0.98 sec
ARIMA(1,1,1)(0,1,0)[12]	: BIC=-4816.196, Time=0.86 sec
ARIMA(1,1,1)(1,1,0)[12]	: BIC=-281.245, Time=1.98 sec
ARIMA(1,1,1)(0,1,1)[12]	: BIC=-503.947, Time=2.16 sec
ARIMA(1,1,1)(1,1,1)[12]	: BIC=2629.702, Time=2.37 sec
ARIMA(0,1,1)(0,1,0)[12]	: BIC=-4821.368, Time=1.94 sec
ARIMA(0,1,1)(1,1,0)[12]	: BIC=-284.334, Time=2.89 sec
ARIMA(0,1,1)(1,1,1)[12]	: BIC=3446.380, Time=2.18 sec
ARIMA(0,1,2)(0,1,0)[12]	: BIC=34414.848, Time=0.81 sec
ARIMA(1,1,2)(0,1,0)[12]	: BIC=49399.697, Time=0.88 sec
ARIMA(0,1,1)(0,1,0)[12] intercept	: BIC=-4816.309, Time=0.86 sec

Best model: ARIMA(0,1,1)(0,1,0)[12]
Total fit time: 35.948 seconds

CV Fold 17/21

Performing stepwise search to minimize bic

ARIMA(0,1,0)(0,1,0)[12]	: BIC=-4793.612, Time=0.41 sec
ARIMA(1,1,0)(1,1,0)[12]	: BIC=-398.294, Time=1.83 sec
ARIMA(0,1,1)(0,1,1)[12]	: BIC=-726.093, Time=2.34 sec
ARIMA(0,1,0)(1,1,0)[12]	: BIC=-437.559, Time=3.59 sec

ARIMA(0,1,0)(0,1,1)[12]	: BIC=-415.619, Time=1.65 sec
ARIMA(0,1,0)(1,1,1)[12]	: BIC=-4550.894, Time=1.69 sec
ARIMA(1,1,0)(0,1,0)[12]	: BIC=-4906.500, Time=0.46 sec
ARIMA(1,1,0)(0,1,1)[12]	: BIC=-724.855, Time=1.79 sec
ARIMA(1,1,0)(1,1,1)[12]	: BIC=-4547.669, Time=1.82 sec
ARIMA(2,1,0)(0,1,0)[12]	: BIC=-1447.662, Time=0.98 sec
ARIMA(1,1,1)(0,1,0)[12]	: BIC=-4898.010, Time=0.84 sec
ARIMA(0,1,1)(0,1,0)[12]	: BIC=-4904.711, Time=1.90 sec
ARIMA(2,1,1)(0,1,0)[12]	: BIC=-3653.476, Time=2.16 sec
ARIMA(1,1,0)(0,1,0)[12] intercept	: BIC=-4901.022, Time=0.50 sec

Best model: ARIMA(1,1,0)(0,1,0)[12]

Total fit time: 21.968 seconds

CV Fold 18/21

Performing stepwise search to minimize bic

ARIMA(0,1,0)(0,1,0)[12]	: BIC=-4793.612, Time=0.38 sec
ARIMA(1,1,0)(1,1,0)[12]	: BIC=-324.307, Time=1.85 sec
ARIMA(0,1,1)(0,1,1)[12]	: BIC=-517.151, Time=2.02 sec
ARIMA(0,1,0)(1,1,0)[12]	: BIC=-461.514, Time=1.62 sec
ARIMA(0,1,0)(0,1,1)[12]	: BIC=-181.793, Time=1.69 sec
ARIMA(0,1,0)(1,1,1)[12]	: BIC=-382.732, Time=2.03 sec
ARIMA(1,1,0)(0,1,0)[12]	: BIC=-4907.407, Time=0.76 sec
ARIMA(1,1,0)(0,1,1)[12]	: BIC=-279.756, Time=3.31 sec
ARIMA(1,1,0)(1,1,1)[12]	: BIC=-501.156, Time=1.95 sec
ARIMA(2,1,0)(0,1,0)[12]	: BIC=-1786.417, Time=0.97 sec
ARIMA(1,1,1)(0,1,0)[12]	: BIC=-4899.646, Time=0.86 sec
ARIMA(0,1,1)(0,1,0)[12]	: BIC=-4906.238, Time=0.80 sec
ARIMA(2,1,1)(0,1,0)[12]	: BIC=-4824.209, Time=0.91 sec
ARIMA(1,1,0)(0,1,0)[12] intercept	: BIC=-4901.890, Time=0.49 sec

Best model: ARIMA(1,1,0)(0,1,0)[12]

Total fit time: 19.671 seconds

CV Fold 19/21

Performing stepwise search to minimize bic

ARIMA(0,1,0)(0,1,0)[12]	: BIC=-4791.576, Time=0.39 sec
ARIMA(1,1,0)(1,1,0)[12]	: BIC=-339.646, Time=1.81 sec
ARIMA(0,1,1)(0,1,1)[12]	: BIC=-216.170, Time=4.21 sec
ARIMA(0,1,0)(1,1,0)[12]	: BIC=-344.894, Time=1.66 sec
ARIMA(0,1,0)(0,1,1)[12]	: BIC=-142.454, Time=1.71 sec
ARIMA(0,1,0)(1,1,1)[12]	: BIC=-463.131, Time=1.82 sec
ARIMA(1,1,0)(0,1,0)[12]	: BIC=-4907.496, Time=0.48 sec
ARIMA(1,1,0)(0,1,1)[12]	: BIC=-174.906, Time=1.81 sec
ARIMA(1,1,0)(1,1,1)[12]	: BIC=-476.694, Time=1.98 sec
ARIMA(2,1,0)(0,1,0)[12]	: BIC=-4343.227, Time=2.26 sec
ARIMA(1,1,1)(0,1,0)[12]	: BIC=15646601.124, Time=1.90 sec
ARIMA(0,1,1)(0,1,0)[12]	: BIC=-3138.793, Time=0.85 sec

```

ARIMA(2,1,1)(0,1,0)[12] : BIC=13733.912, Time=1.07 sec
ARIMA(1,1,0)(0,1,0)[12] intercept : BIC=-4902.007, Time=0.54 sec

```

Best model: ARIMA(1,1,0)(0,1,0)[12]
Total fit time: 22.512 seconds

CV Fold 20/21

Performing stepwise search to minimize bic

```

ARIMA(0,1,0)(0,1,0)[12] : BIC=-856.716, Time=0.37 sec
ARIMA(1,1,0)(1,1,0)[12] : BIC=-1993.879, Time=1.66 sec
ARIMA(0,1,1)(0,1,1)[12] : BIC=-2031.632, Time=1.80 sec
ARIMA(0,1,1)(0,1,0)[12] : BIC=-1920.588, Time=0.79 sec
ARIMA(0,1,1)(1,1,1)[12] : BIC=-2026.077, Time=1.99 sec
ARIMA(0,1,1)(0,1,2)[12] : BIC=-2028.201, Time=5.93 sec
ARIMA(0,1,1)(1,1,0)[12] : BIC=-2000.510, Time=1.55 sec
ARIMA(0,1,1)(1,1,2)[12] : BIC=-2021.391, Time=4.23 sec
ARIMA(0,1,0)(0,1,1)[12] : BIC=-2071.689, Time=1.48 sec
ARIMA(0,1,0)(1,1,1)[12] : BIC=-2061.716, Time=3.80 sec
ARIMA(0,1,0)(0,1,2)[12] : BIC=-2066.135, Time=3.24 sec
ARIMA(0,1,0)(1,1,0)[12] : BIC=-2016.711, Time=1.44 sec
ARIMA(0,1,0)(1,1,2)[12] : BIC=-2059.875, Time=3.58 sec
ARIMA(1,1,0)(0,1,1)[12] : BIC=-2021.063, Time=1.67 sec
ARIMA(1,1,1)(0,1,1)[12] : BIC=-2026.165, Time=3.98 sec
ARIMA(0,1,0)(0,1,1)[12] intercept : BIC=-2022.221, Time=1.62 sec

```

Best model: ARIMA(0,1,0)(0,1,1)[12]
Total fit time: 39.146 seconds

CV Fold 21/21

Performing stepwise search to minimize bic

```

ARIMA(0,1,0)(0,1,0)[12] : BIC=730.121, Time=0.39 sec
ARIMA(1,1,0)(1,1,0)[12] : BIC=-389.685, Time=1.66 sec
ARIMA(0,1,1)(0,1,1)[12] : BIC=-413.192, Time=1.82 sec
ARIMA(0,1,1)(0,1,0)[12] : BIC=-335.791, Time=0.77 sec
ARIMA(0,1,1)(1,1,1)[12] : BIC=-400.474, Time=2.01 sec
ARIMA(0,1,1)(0,1,2)[12] : BIC=-408.908, Time=5.92 sec
ARIMA(0,1,1)(1,1,0)[12] : BIC=-389.263, Time=1.55 sec
ARIMA(0,1,1)(1,1,2)[12] : BIC=-402.672, Time=4.21 sec
ARIMA(0,1,0)(0,1,1)[12] : BIC=-435.284, Time=1.48 sec
ARIMA(0,1,0)(1,1,1)[12] : BIC=-423.492, Time=3.35 sec
ARIMA(0,1,0)(0,1,2)[12] : BIC=-419.419, Time=3.67 sec
ARIMA(0,1,0)(1,1,0)[12] : BIC=-404.997, Time=1.48 sec
ARIMA(0,1,0)(1,1,2)[12] : BIC=-426.285, Time=3.53 sec
ARIMA(1,1,0)(0,1,1)[12] : BIC=-413.848, Time=1.60 sec
ARIMA(1,1,1)(0,1,1)[12] : BIC=-407.454, Time=3.86 sec
ARIMA(0,1,0)(0,1,1)[12] intercept : BIC=-412.642, Time=1.70 sec

```

Best model: ARIMA(0,1,0)(0,1,1)[12]

Total fit time: 39.049 seconds

Final Cross-Validation Accuracy: 0.0000

Performing stepwise search to minimize bic

ARIMA(0,1,0)(0,1,0)[12]	: BIC=8654.887, Time=3.08 sec
ARIMA(1,1,0)(1,1,0)[12]	: BIC=49.999, Time=16.63 sec
ARIMA(0,1,1)(0,1,1)[12]	: BIC=-153.506, Time=20.32 sec
ARIMA(0,1,1)(0,1,0)[12]	: BIC=453.176, Time=9.80 sec
ARIMA(0,1,1)(1,1,1)[12]	: BIC=-103.822, Time=21.04 sec
ARIMA(0,1,1)(0,1,2)[12]	: BIC=-149.619, Time=36.59 sec
ARIMA(0,1,1)(1,1,0)[12]	: BIC=51.931, Time=16.36 sec
ARIMA(0,1,1)(1,1,2)[12]	: BIC=-154.199, Time=37.48 sec
ARIMA(0,1,1)(2,1,2)[12]	: BIC=-133.330, Time=40.32 sec
ARIMA(0,1,1)(2,1,1)[12]	: BIC=-78.776, Time=33.65 sec
ARIMA(0,1,0)(1,1,2)[12]	: BIC=-395.808, Time=32.63 sec
ARIMA(0,1,0)(0,1,2)[12]	: BIC=-363.581, Time=29.57 sec
ARIMA(0,1,0)(1,1,1)[12]	: BIC=-293.575, Time=18.93 sec
ARIMA(0,1,0)(2,1,2)[12]	: BIC=-364.285, Time=34.61 sec
ARIMA(0,1,0)(0,1,1)[12]	: BIC=-362.949, Time=17.75 sec
ARIMA(0,1,0)(2,1,1)[12]	: BIC=-257.973, Time=32.04 sec
ARIMA(1,1,0)(1,1,2)[12]	: BIC=-156.177, Time=36.71 sec
ARIMA(1,1,1)(1,1,2)[12]	: BIC=-149.429, Time=40.65 sec
ARIMA(0,1,0)(1,1,2)[12] intercept	: BIC=-157.122, Time=34.61 sec

Best model: ARIMA(0,1,0)(1,1,2)[12]


Total fit time: 512.810 seconds

GARCH model fitted. AIC: -399.12

MAE: 93.3231

Directional Accuracy: 0.7595

Running model diagnostics...



Ljung-Box p-value: 1.0000
 Jarque-Bera p-value: 0.0000

3.2.2 ENSEMBLE MODEL

```
[131]: test_actual_prices = test_df['Close'].values
lr_pred_prices = np.exp(y_pred.flatten())
rf_pred_prices = np.exp(y_pred_rf)
arima_pred_prices = test_df['Forecast_Close'].values
min_len = min(len(test_actual_prices), len(lr_pred_prices),
               len(rf_pred_prices), len(arima_pred_prices))
test_actual_prices = test_actual_prices[:min_len]
lr_pred_prices = lr_pred_prices[:min_len]
```

```

rf_pred_prices = rf_pred_prices[:min_len]
arima_pred_prices = arima_pred_prices[:min_len]
weights = np.array([1/mae_lr, 1/mae_rf, 1/mae_ar])
normalized_weights = weights / weights.sum()
print(f"\n[Ensemble] Model Weights:")
print(f"Linear Regression: {normalized_weights[0]:.2%}")
print(f"Random Forest: {normalized_weights[1]:.2%}")
print(f"ARIMA: {normalized_weights[2]:.2%}")
ensemble_predictions = (
    normalized_weights[0] * lr_pred_prices +
    normalized_weights[1] * rf_pred_prices +
    normalized_weights[2] * arima_pred_prices
)
mae_ensemble = mean_absolute_error(test_actual_prices, ensemble_predictions)
actual_directions = np.sign(np.diff(test_actual_prices)).astype(int)
ensemble_directions = np.sign(ensemble_predictions[1:] - test_actual_prices[:
    ↪-1]).astype(int)
non_zero_mask = actual_directions != 0
accuracy_ensemble = (ensemble_directions[non_zero_mask] ==
    ↪actual_directions[non_zero_mask]).mean()
print(f"\n[Ensemble] Performance:")
print(f"MAE: {mae_ensemble:.4f}")
print(f"Directional Accuracy: {accuracy_ensemble:.2%}")
plt.figure(figsize=(14, 7))
plt.plot(test_df.index[:min_len], test_actual_prices,
    label='Actual Prices', color='black', alpha=0.6)
plt.plot(test_df.index[:min_len], ensemble_predictions,
    label='Ensemble Predictions', color='red', linestyle='--')
plt.title('Ensemble Model vs Actual Prices (AMZN)')
plt.xlabel('Date')
plt.ylabel('Price ($)')
plt.legend()
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

```

```


[Ensemble] Model Weights:
Linear Regression: 92.22%
Random Forest: 7.77%
ARIMA: 0.01%

```

```

[Ensemble] Performance:
MAE: 175.7067
Directional Accuracy: 45.76%

```



3.2.3 TEMPORAL FUSION TRANSFORMER

```
[132]: import torch
from pytorch_forecasting import TemporalFusionTransformer, TimeSeriesDataSet
from pytorch_forecasting.data import GroupNormalizer
from pytorch_forecasting.metrics import QuantileLoss
from lightning.pytorch.callbacks import EarlyStopping, LearningRateMonitor
from lightning.pytorch.loggers import TensorBoardLogger
import lightning.pytorch as pl
SEED = 42
os.environ['PYTHONHASHSEED'] = str(SEED)
random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed(SEED)
torch.cuda.manual_seed_all(SEED)
torch.backends.cudnn.deterministic = True
torch.backends.cudnn.benchmark = False
os.environ['CUBLAS_WORKSPACE_CONFIG'] = ':4096:8'
torch.use_deterministic_algorithms(True, warn_only=True)
amzn_df_tft = filtered_df_sorted.loc['amzn'].copy().reset_index()
amzn_df_tft['Ticker'] = 'amzn'
amzn_df_tft = amzn_df_tft[amzn_df_tft['Close'] > 0]
amzn_df_tft['Close_Log'] = np.log(amzn_df_tft['Close'])
numerical_cols = [
    'Open', 'High', 'Low', 'Close', 'Volume',
    'Daily Return', '7-Day Moving Average', '30-Day Moving Average',
    'Rolling Volatility (30d)', 'Price Change', 'High-Low Spread',
```

```

    'Intraday Volatility %', 'RSI_14', 'MACD',
    'Volume_Change', 'Volume_Price_Ratio'
]
for col in numerical_cols:
    if col in amzn_df_tft.columns:
        amzn_df_tft[col] = pd.to_numeric(amzn_df_tft[col], errors='coerce')
        amzn_df_tft[col] = amzn_df_tft[col].replace([np.inf, -np.inf], np.nan)
        amzn_df_tft[col] = amzn_df_tft[col].fillna(method='ffill').
        ↪fillna(method='bfill')
        amzn_df_tft[col] = amzn_df_tft[col].astype(np.float32)
amzn_df_tft = amzn_df_tft.sort_values(['Ticker', 'Date'])
amzn_df_tft['time_idx'] = amzn_df_tft.groupby('Ticker').cumcount().astype(int)
date_features = ['Year', 'Month', 'Day_of_Week', 'Number_of_Days']
for col in date_features:
    if col in amzn_df_tft.columns:
        amzn_df_tft[col] = amzn_df_tft[col].astype(int)
categorical_features = ['Ticker']
time_varying_known_reals = [
    'time_idx', 'Year', 'Month', 'Day_of_Week', 'Number_of_Days',
    '7-Day Moving Average', '30-Day Moving Average'
]
time_varying_unknown_reals = [
    'Open', 'High', 'Low', 'Close', 'Volume', 'Daily Return',
    'Rolling Volatility (30d)', 'Price Change', 'High-Low Spread',
    'Intraday Volatility %', 'RSI_14', 'MACD', 'Volume_Change', ↵
    ↪'Volume_Price_Ratio'
]
target = 'Close_Log'
max_prediction_length = 20
max_encoder_length = 60
training_cutoff = amzn_df_tft['time_idx'].max() - max_prediction_length
cols_for_dataset = [target] + categorical_features + time_varying_known_reals + ↵
    ↪time_varying_unknown_reals
tft_data = amzn_df_tft[cols_for_dataset].copy()
for col in time_varying_known_reals + time_varying_unknown_reals + [target]:
    tft_data[col] = tft_data[col].astype(np.float32)
tft_data["time_idx"] = tft_data["time_idx"].astype(int)
tft_data['Ticker'] = tft_data['Ticker'].astype(str)
print("\nData types in final dataset:")
print(tft_data.dtypes)
training_data = TimeSeriesDataSet(
    tft_data[tft_data['time_idx'] <= training_cutoff],
    time_idx="time_idx",
    target=target,
    group_ids=["Ticker"],
    min_encoder_length=max_encoder_length // 2,
    max_encoder_length=max_encoder_length,

```

```

min_prediction_length=1,
max_prediction_length=max_prediction_length,
static_categoricals=categorical_features,
time_varying_known_reals=time_varying_known_reals,
time_varying_unknown_reals=time_varying_unknown_reals,
target_normalizer=GroupNormalizer(
    groups=["Ticker"],
    transformation="softplus",
    center=False
),
add_relative_time_idx=True,
add_target_scales=True,
add_encoder_length=True,
allow_missing_timesteps=False
)
validation_data = TimeSeriesDataSet.from_dataset(
    training_data,
    tft_data,
    predict=True,
    stop_randomization=True,
)
batch_size = 64
train_dataloader = training_data.to_dataloader(
    train=True,
    batch_size=batch_size,
    num_workers=0,
    shuffle=False,
    drop_last=False,
    pin_memory=False
)
val_dataloader = validation_data.to_dataloader(
    train=False,
    batch_size=batch_size*4,
    num_workers=0,
    shuffle=False,
    drop_last=False,
    pin_memory=False
)
tft = TemporalFusionTransformer.from_dataset(
    training_data,
    hidden_size=32,
    dropout=0.1,
    hidden_continuous_size=16,
    attention_head_size=2,
    learning_rate=0.01,
    loss=QuantileLoss(quantiles=[0.05, 0.5, 0.95]),
    optimizer="adamw",

```



```

        reduce_on_plateau_patience=3,
    )
    early_stop_callback = EarlyStopping(
        monitor="val_loss",
        min_delta=1e-4,
        patience=10,
        verbose=True,
        mode="min"
    )
    lr_logger = LearningRateMonitor()
    logger = TensorBoardLogger("lightning_logs", name="tft_amzn")
    trainer = pl.Trainer(
        max_epochs=20,
        accelerator="auto",
        devices="auto",
        callbacks=[lr_logger, early_stop_callback],
        logger=logger,
        gradient_clip_val=0.1,
        enable_progress_bar=True,
        deterministic=True,
        benchmark=False,
        enable_model_summary=True,
        precision="32-true",
    )
    trainer.fit(
        tft,
        train_dataloaders=train_dataloader,
        val_dataloaders=val_dataloader
    )
    tft.eval()
    with torch.no_grad(), torch.inference_mode():
        predictions = tft.predict(val_dataloader)
        actuals = torch.cat([y[0] for x, y in iter(val_dataloader)])
    actuals_orig = torch.exp(actuals)
    predictions_orig = torch.exp(predictions)
    mae_orig = (actuals_orig - predictions_orig).abs().mean()
    print(f"Original Scale MAE: {mae_orig:.2f}")
    actual_prices = actuals_orig.cpu().numpy().flatten()
    predicted_prices = predictions_orig.cpu().numpy().flatten()
    actual_prev = actual_prices[:-1]
    actual_curr = actual_prices[1:]
    predicted_prev = predicted_prices[:-1]
    predicted_curr = predicted_prices[1:]
    actual_direction = np.sign(actual_curr - actual_prev)
    predicted_direction = np.sign(predicted_curr - predicted_prev)
    direction_matches = actual_direction == predicted_direction
    accuracy_tft = np.mean(direction_matches)

```

```
print(f"Direction Accuracy: {accuracy_tft * 100:.2f}%")
best_tft_model = tft
```

Data types in final dataset:

```
Close_Log          float32
Ticker             object
time_idx           int64
Year               float32
Month              float32
Day_of_Week        float32
Number_of_Days     float32
7-Day Moving Average float32
30-Day Moving Average float32
Open               float32
High               float32
Low                float32
Close              float32
Volume             float32
Daily Return       float32
Rolling Volatility (30d) float32
Price Change       float32
High-Low Spread    float32
Intraday Volatility % float32
RSI_14             float32
MACD               float32
Volume_Change      float32
Volume_Price_Ratio float32
dtype: object
```

INFO: GPU available: False, used: False

INFO:lightning.pytorch.utilities.rank_zero:GPU available: False, used: False

INFO: TPU available: False, using: 0 TPU cores

INFO:lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU cores

INFO: IPU available: False, using: 0 IPUs

INFO:lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPUs

INFO: HPU available: False, using: 0 HPUs

INFO:lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPUs

INFO:

	Name	Type	
Params			

0	loss	QuantileLoss	0
1	logging_metrics	ModuleList	0
2	input_embeddings	MultiEmbedding	1
3	prescalers	ModuleDict	800

4		static_variable_selection		VariableSelectionNetwork		5.8
K						
5		encoder_variable_selection		VariableSelectionNetwork		50.2
K						
6		decoder_variable_selection		VariableSelectionNetwork		16.1
K						
7		static_context_variable_selection		GatedResidualNetwork		4.3
K						
8		static_context_initial_hidden_lstm		GatedResidualNetwork		4.3
K						
9		static_context_initial_cell_lstm		GatedResidualNetwork		4.3
K						
10		static_context_enrichment		GatedResidualNetwork		4.3
K						
11		lstm_encoder		LSTM		8.4
K						
12		lstm_decoder		LSTM		8.4
K						
13		post_lstm_gate_encoder		GatedLinearUnit		2.1
K						
14		post_lstm_add_norm_encoder		AddNorm		64
15		static_enrichment		GatedResidualNetwork		5.3
K						
16		multihead_attn		InterpretableMultiHeadAttention		3.2
K						
17		post_attn_gate_norm		GateAddNorm		2.2
K						
18		pos_wise_ff		GatedResidualNetwork		4.3
K						
19		pre_output_gate_norm		GateAddNorm		2.2
K						
20		output_layer		Linear		99

```

-----
125 K      Trainable params
0          Non-trainable params
125 K      Total params
0.502      Total estimated model params size (MB)
INFO:lightning.pytorch.callbacks.model_summary:

```

Name	Type	
------	------	--

0		loss		QuantileLoss		0
1		logging_metrics		ModuleList		0
2		input_embeddings		MultiEmbedding		1
3		prescalers		ModuleDict		800
4		static_variable_selection		VariableSelectionNetwork		5.8

K			
5	encoder_variable_selection	VariableSelectionNetwork	50.2
K			
6	decoder_variable_selection	VariableSelectionNetwork	16.1
K			
7	static_context_variable_selection	GatedResidualNetwork	4.3
K			
8	static_context_initial_hidden_lstm	GatedResidualNetwork	4.3
K			
9	static_context_initial_cell_lstm	GatedResidualNetwork	4.3
K			
10	static_context_enrichment	GatedResidualNetwork	4.3
K			
11	lstm_encoder	LSTM	8.4
K			
12	lstm_decoder	LSTM	8.4
K			
13	post_lstm_gate_encoder	GatedLinearUnit	2.1
K			
14	post_lstm_add_norm_encoder	AddNorm	64
15	static_enrichment	GatedResidualNetwork	5.3
K			
16	multihead_attn	InterpretableMultiHeadAttention	3.2
K			
17	post_attn_gate_norm	GateAddNorm	2.2
K			
18	pos_wise_ff	GatedResidualNetwork	4.3
K			
19	pre_output_gate_norm	GateAddNorm	2.2
K			
20	output_layer	Linear	99

```

-----
125 K      Trainable params
0          Non-trainable params
125 K      Total params
0.502      Total estimated model params size (MB)

```

Sanity Checking: 0it [00:00, ?it/s]

Training: 0it [00:00, ?it/s]

Validation: 0it [00:00, ?it/s]

INFO: Metric val_loss improved. New best score: 0.351

INFO:lightning.pytorch.callbacks.early_stopping:Metric val_loss improved. New best score: 0.351

Validation: 0it [00:00, ?it/s]

INFO: Metric val_loss improved by 0.106 >= min_delta = 0.0001. New best score:

```
0.245
INFO:lightning.pytorch.callbacks.early_stopping:Metric val_loss improved by
0.106 >= min_delta = 0.0001. New best score: 0.245
Validation: 0it [00:00, ?it/s]
Validation: 0it [00:00, ?it/s]
INFO: Metric val_loss improved by 0.025 >= min_delta = 0.0001. New best score:
0.220
INFO:lightning.pytorch.callbacks.early_stopping:Metric val_loss improved by
0.025 >= min_delta = 0.0001. New best score: 0.220
Validation: 0it [00:00, ?it/s]
INFO: Metric val_loss improved by 0.027 >= min_delta = 0.0001. New best score:
0.193
INFO:lightning.pytorch.callbacks.early_stopping:Metric val_loss improved by
0.027 >= min_delta = 0.0001. New best score: 0.193
Validation: 0it [00:00, ?it/s]
INFO: Metric val_loss improved by 0.084 >= min_delta = 0.0001. New best score:
0.110
INFO:lightning.pytorch.callbacks.early_stopping:Metric val_loss improved by
0.084 >= min_delta = 0.0001. New best score: 0.110
Validation: 0it [00:00, ?it/s]
Validation: 0it [00:00, ?it/s]
INFO: Metric val_loss improved by 0.078 >= min_delta = 0.0001. New best score:
0.031
INFO:lightning.pytorch.callbacks.early_stopping:Metric val_loss improved by
0.078 >= min_delta = 0.0001. New best score: 0.031
Validation: 0it [00:00, ?it/s]
Validation: 0it [00:00, ?it/s]
Validation: 0it [00:00, ?it/s]
Validation: 0it [00:00, ?it/s]
Validation: 0it [00:00, ?it/s]
Validation: 0it [00:00, ?it/s]
Validation: 0it [00:00, ?it/s]
Validation: 0it [00:00, ?it/s]
Validation: 0it [00:00, ?it/s]
Validation: 0it [00:00, ?it/s]
INFO: Metric val_loss improved by 0.011 >= min_delta = 0.0001. New best score:
0.021
```

```

INFO:lightning.pytorch.callbacks.early_stopping:Metric val_loss improved by
0.011 >= min_delta = 0.0001. New best score: 0.021

Validation: 0it [00:00, ?it/s]

Validation: 0it [00:00, ?it/s]

INFO: `Trainer.fit` stopped: `max_epochs=20` reached.
INFO:lightning.pytorch.utilities.rank_zero:`Trainer.fit` stopped:
`max_epochs=20` reached.
INFO: GPU available: False, used: False
INFO:lightning.pytorch.utilities.rank_zero:GPU available: False, used: False
INFO: TPU available: False, using: 0 TPU cores
INFO:lightning.pytorch.utilities.rank_zero:TPU available: False, using: 0 TPU
cores
INFO: IPU available: False, using: 0 IPUs
INFO:lightning.pytorch.utilities.rank_zero:IPU available: False, using: 0 IPUs
INFO: HPU available: False, using: 0 HPUs
INFO:lightning.pytorch.utilities.rank_zero:HPU available: False, using: 0 HPUs

Original Scale MAE: 37.58
Direction Accuracy: 52.63%

```


3.3 MODEL INTERPRETABILITY

I visualized the importance of features from the Random Forest model and analyzed the SARIMAX model components and the impact of exogenous variables.

```

[133]: if 'rf_model' in locals():
        feature_importance = pd.Series(rf_model.feature_importances_,
        ↪index=X_train_rf.columns)
        feature_importance.sort_values(ascending=False, inplace=True)
        plt.figure(figsize=(12, 8))
        sns.barplot(x=feature_importance.values, y=feature_importance.index)
        plt.title('Random Forest Feature Importance')
        plt.xlabel('Importance')
        plt.ylabel('Feature')
        plt.tight_layout()
        plt.show()
        print("\nTop 10 Random Forest Features by Importance:")
        print(feature_importance.head(10))
    else:
        print("Random Forest model (rf_model) not found. Please ensure the Random_
        ↪Forest training section ran successfully.")

```



Top 10 Random Forest Features by Importance:

Low	0.638423
High	0.286339
Open	0.055518
7-Day Moving Average	0.013302
Number_of_Days	0.005750
30-Day Moving Average	0.000604
MACD	0.000017
Rolling Volatility (30d)	0.000013
Intraday Volatility %	0.000009
High-Low Spread	0.000008

dtype: float64

```
[134]: from statsmodels.stats.diagnostic import acorr_ljungbox
from statsmodels.stats.diagnostic import het_arch
try:
    if 'sarimax_model' in locals():
        if hasattr(sarimax_model, 'arima_res_'):
            results = sarimax_model.arima_res_
            print("Using pmdarima results wrapper")
        elif hasattr(sarimax_model, 'results'):
            results = sarimax_model.results
```

```

        print("Using statsmodels results wrapper")
    else:
        results = sarimax_model
        print("Using base model as results")
    print("\nSARIMAX Model Summary:")
    print(results.summary())
    residuals = results.resid.dropna()
    print(f"\nResiduals analysis (n={len(residuals)}):")
    print(f"Mean: {residuals.mean():.6f}, Std: {residuals.std():.6f}")
    fig = plt.figure(figsize=(15, 18))
    gs = fig.add_gridspec(4, 1)
    ax1 = fig.add_subplot(gs[0])
    ax1.plot(residuals)
    ax1.set_title('Residuals over Time')
    ax1.grid(alpha=0.3)
    ax2 = fig.add_subplot(gs[1])
    sns.histplot(residuals, kde=True, ax=ax2)
    ax2.set_title('Residual Distribution')
    ax2.annotate(f"Skew: {residuals.skew():.2f}, Kurtosis: {residuals.
↪kurtosis():.2f}",
                xy=(0.7, 0.9), xycoords='axes fraction')
    ax3 = fig.add_subplot(gs[2])
    plot_acf(residuals, lags=40, ax=ax3, title='Residual ACF')
    ax4 = fig.add_subplot(gs[3])
    plot_pacf(residuals, lags=40, ax=ax4, title='Residual PACF')
    plt.tight_layout()
    plt.show()
    try:
        arch_test = het_arch(residuals)
        print(f"\nARCH Test for Volatility Clustering:
↪p-value={arch_test[1]:.4f}")
        if arch_test[1] < 0.05:
            print("Significant ARCH effects detected - volatility modeling
↪needed")
            print("Recommendation: Implement GARCH model for volatility
↪clustering")
        else:
            print("No significant ARCH effects detected")
    except Exception as e:
        print(f"ARCH test failed: {str(e)}")
    lb_test = acorr_ljungbox(residuals, lags=[10], return_df=True)
    print(f"\nLjung-Box Test (lag 10): p-value={lb_test['lb_pvalue'].
↪iloc[0]:.4f}")
    if lb_test['lb_pvalue'].iloc[0] < 0.05:
        print("Significant autocorrelation detected - model improvement
↪needed")

```



```

        else:
            print("No significant autocorrelation detected")
        from scipy.stats import jarque_bera
        jb_test = jarque_bera(residuals)
        print(f"\nJarque-Bera Test: p-value={jb_test[1]:.4f}")
        if jb_test[1] < 0.05:
            print("Residuals are not normally distributed")
        else:
            print("Residuals are normally distributed")
    else:
        print("SARIMAX model not found in memory")
except Exception as e:
    print(f"\nError in diagnostics: {str(e)}")
    if 'sarimax_model' in locals() and hasattr(sarimax_model, 'summary'):
        print("\nBasic Model Summary:")
        print(sarimax_model.summary())

```

SARIMAX model not found in memory

3.4 EVALUATION

Using Mean Absolute Error (MAE) and direction accuracy as evaluation metrics of the models, I evaluated and compared the SARIMAX, Ensemble Model, TFT and the other models I tried and tested in the second week of this project.

```

[135]: mae_results = {
        'Linear Regression': mae_lr,
        'Random Forest': mae_rf,
        'ARIMA': mae_ar,
        'SARIMAX': mae_sarimax,
        'Ensemble': mae_ensemble,
        'TFT': mae_orig.item()
    }
    accuracy_results = {
        'Linear Regression': accuracy_lr,
        'Random Forest': accuracy_rf,
        'ARIMA': accuracy_ar,
        'SARIMAX': accuracy_sarimax,
        'Ensemble': accuracy_ensemble,
        'TFT': accuracy_tft
    }
    print("{:<20} | {:<10} | {:<18}".format("Model", "MAE", "Direction Accuracy"))
    print("-" * 50)
    for model_name in mae_results:
        print("{:<20} | {:<10.4f} | {:<18.4f}".format(
            model_name,
            mae_results[model_name],
            accuracy_results[model_name]

```

```

    ))
best_mae_model_name = min(mae_results, key=mae_results.get)
print(f"\nModel with the lowest MAE: {best_mae_model_name} (MAE:␣
    ↳{mae_results[best_mae_model_name]:.4f})")
valid_accuracy_results = {k: v for k, v in accuracy_results.items() if v > 0}
if valid_accuracy_results:
    best_accuracy_model_name = max(valid_accuracy_results,␣
    ↳key=valid_accuracy_results.get)
    print(f"Model with the highest Direction Accuracy:␣
    ↳{best_accuracy_model_name} (Accuracy:␣
    ↳{accuracy_results[best_accuracy_model_name]:.4f})")
    best_overall_model_name = best_accuracy_model_name
else:
    print("Could not determine a best model based on accuracy as all have 0␣
    ↳accuracy.")
    best_overall_model_name = best_mae_model_name
print(f"\nBased on direction accuracy, the best performing model is:␣
    ↳{best_overall_model_name}")

```

Model	MAE	Direction Accuracy
Linear Regression	0.0140	0.5499
Random Forest	0.1660	0.6106
ARIMA	138.7132	0.4676
SARIMAX	93.3231	0.7595
Ensemble	175.7067	0.4576
TFT	37.5778	0.5263

Model with the lowest MAE: Linear Regression (MAE: 0.0140)

Model with the highest Direction Accuracy: SARIMAX (Accuracy: 0.7595)

Based on direction accuracy, the best performing model is: SARIMAX

3.5 BACKTESTING AND RISK MANAGEMENT

I backtested SARIMAX model using the same strategy I used in the second week, enhancing it with two more features, i.e., Stop-loss (exit if a price drops 5% from entry) and Sharpe Ratio.

```

[136]: from datetime import datetime
train_df.to_csv('train_df.csv')
test_df.to_csv('test_df.csv')
amzn_df.to_csv('amzn_df.csv')
df = pd.read_csv('test_df.csv', parse_dates=['Date'])
df.sort_values('Date', inplace=True)
df['Predicted_Direction'] = df['Predicted_Direction'].astype(bool)
pv = 1.0

```

```


shares = 0
in_market = False
portfolio_values = [pv]
daily_returns = []
for i in range(len(df)):
    row = df.iloc[i]
    open_price = row['Open']
    close_price = row['Close']
    pred = row['Predicted_Direction']
    if not in_market and pred:
        shares = pv / open_price
        in_market = True
    elif in_market and not pred:
        pv = shares * open_price
        shares = 0
        in_market = False
    pv_end = shares * close_price if in_market else pv
    daily_return = (pv_end - portfolio_values[-1]) / portfolio_values[-1]
    daily_returns.append(daily_return)
    portfolio_values.append(pv_end)
    pv = pv_end
strategy_profit = portfolio_values[-1] - 1.0
buy_hold_return = (df['Close'].iloc[-1] / df['Open'].iloc[0]) - 1
daily_returns = np.array(daily_returns)
if daily_returns.std() == 0:
    sharpe_ratio = np.nan
else:
    sharpe_ratio = (daily_returns.mean() / daily_returns.std()) * np.sqrt(252)
start_date = df['Date'].iloc[0]
end_date = df['Date'].iloc[-1]
duration_days = (end_date - start_date).days
if duration_days > 0:
    cagr = (portfolio_values[-1] ** (365.25/duration_days)) - 1
else:
    cagr = np.nan
bh_values = [1.0]
for i in range(len(df)):
    if i == 0:
        bh_return = df['Close'].iloc[i] / df['Open'].iloc[0]
    else:
        bh_return = df['Close'].iloc[i] / df['Close'].iloc[i-1]
    bh_values.append(bh_values[-1] * bh_return)
plt.figure(figsize=(12, 6))
plt.plot(df['Date'], portfolio_values[1:], label='Trading Strategy')
plt.plot(df['Date'], bh_values[1:], label='Buy and Hold')
plt.xlabel('Date')
plt.ylabel('Portfolio Value')

```

```

plt.title('Strategy vs Buy-and-Hold Performance')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
print(f"{'Strategy Profit': '<25'} {strategy_profit:.4f} ({strategy_profit*100:.2f}%)")
print(f"{'Buy and Hold Profit': '<25'} {buy_hold_return:.4f} ({buy_hold_return*100:.2f}%)")
print(f"{'Sharpe Ratio': '<25'} {sharpe_ratio:.4f}")
print(f"{'Annual Return': '<25'} {cagr*100:.2f}%")
print(f"{'Start Date': '<25'} {start_date.date()}")
print(f"{'End Date': '<25'} {end_date.date()}")

```



Strategy Profit:	1.3143 (131.43%)
Buy and Hold Profit:	0.5430 (54.30%)
Sharpe Ratio:	3.7825
Annual Return:	128.44%
Start Date:	2015-04-30
End Date:	2016-05-05

3.6 PORTFOLIO OPTIMIZATION

Using historical daily returns of AAPL and AMZN, I calculated the expected returns and the covariance matrix, then performed Mean-Variance Optimization (MVO) to determine the optimal portfolio allocation. The Sharpe Ratio was used as the objective function to maximize risk-adjusted returns, resulting in the optimal asset weights.

```

[137]: from scipy.optimize import minimize
tickers = ['amzn', 'aapl']
portfolio_df = filtered_df_sorted.loc[tickers, 'Close'].unstack(level=0).copy()
portfolio_df = portfolio_df.dropna()
portfolio_df = portfolio_df.sort_index()
returns = portfolio_df.pct_change().dropna()
mean_returns = returns.mean()
cov_matrix = returns.cov()
def portfolio_performance(weights, mean_returns, cov_matrix, risk_free_rate=0.
    ↪01/252):
    returns = np.dot(weights, mean_returns)
    std = np.sqrt(np.dot(weights.T, np.dot(cov_matrix, weights)))
    sharpe = (returns - risk_free_rate) / std
    return -sharpe
num_assets = len(mean_returns)
constraints = ({'type': 'eq', 'fun': lambda x: np.sum(x) - 1})
bounds = tuple((0, 1) for _ in range(num_assets))
initial_guess = num_assets * [1. / num_assets]
opt = minimize(
    portfolio_performance,
    initial_guess,
    args=(mean_returns, cov_matrix),
    method='SLSQP',
    bounds=bounds,
    constraints=constraints
)
optimal_weights = opt.x
print("Optimal Weights (Sharpe-maximizing):")
for ticker, weight in zip(tickers, optimal_weights):
    print(f"{ticker.upper(): <10s}: {weight:.2%}")
sharpe_max = -opt.fun
expected_return = np.dot(optimal_weights, mean_returns)
volatility = np.sqrt(np.dot(optimal_weights.T, np.dot(cov_matrix,
    ↪optimal_weights)))
print(f"\nExpected Return: {expected_return:.4f}")
print(f"Volatility: {volatility:.4f}")
print(f"Sharpe Ratio: {sharpe_max:.4f}")
(returns @ optimal_weights).cumsum().plot(label='Optimized Portfolio',
    ↪figsize=(12,6))
returns['amzn'].cumsum().plot(label='AMZN')
returns['aapl'].cumsum().plot(label='AAPL')
plt.title("Cumulative Returns")
plt.ylabel("Cumulative Return")
plt.legend()
plt.grid(True)
plt.show()

```

Optimal Weights (Sharpe-maximizing):


AMZN: 55.02%

AAPL: 44.98%

Expected Return: 0.0012

Volatility: 0.0204

Sharpe Ratio: 0.0553



4 WEEK 4

I chose to shift my focus from pure point forecasting to classification and uncertainty modelling, using Support Vector Machines (SVM) for directional prediction and Bayesian Regression for confidence-aware price prediction, aiming to combine classification and probabilistic reasoning to make smarter, risk-aware financial decisions on the stock AMZN.

4.1 SUPPORT VECTOR MACHINE

4.1.1 LABEL CREATION

I created a binary label, i.e. 1 if the next day's return is > 0 , else 0.

```
[138]: filtered_df_sorted['label'] = np.where(
        filtered_df_sorted['Daily Return'].notna(),
        (filtered_df_sorted['Daily Return'] > 0).astype(int),
        np.nan
    )
```

4.1.2 FEATURE ENGINEERING

I added more features to the dataset, like Daily returns, Moving averages, RSI, Bollinger Bands, Volume, e.t.c. Using the data from the past 5-10 days.

```
[139]: filtered_df_sorted = filtered_df_sorted.sort_index(level=['Ticker', 'Date'])
filtered_df_sorted['MA_5'] = filtered_df_sorted.groupby('Ticker')['Close'].
    ↪transform(lambda x: x.rolling(5).mean())
filtered_df_sorted['MA_10'] = filtered_df_sorted.groupby('Ticker')['Close'].
    ↪transform(lambda x: x.rolling(10).mean())
def compute_rsi(series, window):
    delta = series.diff()
    gain = delta.clip(lower=0)
    loss = -delta.clip(upper=0)
    avg_gain = gain.rolling(window).mean()
    avg_loss = loss.rolling(window).mean()
    rs = avg_gain / avg_loss
    rsi = 100 - (100 / (1 + rs))
    return rsi
filtered_df_sorted['RSI_5'] = filtered_df_sorted.groupby('Ticker')['Close'].
    ↪transform(lambda x: compute_rsi(x, 5))
filtered_df_sorted['RSI_10'] = filtered_df_sorted.groupby('Ticker')['Close'].
    ↪transform(lambda x: compute_rsi(x, 10))
for window in [5, 10]:
    ma = filtered_df_sorted.groupby('Ticker')['Close'].transform(lambda x: x.
    ↪rolling(window).mean())
    std = filtered_df_sorted.groupby('Ticker')['Close'].transform(lambda x: x.
    ↪rolling(window).std())
    filtered_df_sorted[f'BB_Mid_{window}'] = ma
    filtered_df_sorted[f'BB_Upper_{window}'] = ma + 2 * std
    filtered_df_sorted[f'BB_Lower_{window}'] = ma - 2 * std
filtered_df_sorted['Vol_MA_5'] = filtered_df_sorted.groupby('Ticker')['Volume'].
    ↪transform(lambda x: x.rolling(5).mean())
filtered_df_sorted['Vol_MA_10'] = filtered_df_sorted.
    ↪groupby('Ticker')['Volume'].transform(lambda x: x.rolling(10).mean())
if 'Daily Return' not in filtered_df_sorted.columns:
    filtered_df_sorted['Daily Return'] = filtered_df_sorted.
    ↪groupby('Ticker')['Close'].pct_change()
filtered_df_sorted['Volatility_5'] = filtered_df_sorted.
    ↪groupby('Ticker')['Daily Return'].transform(lambda x: x.rolling(5).std())
filtered_df_sorted['Volatility_10'] = filtered_df_sorted.
    ↪groupby('Ticker')['Daily Return'].transform(lambda x: x.rolling(10).std())
filtered_df_sorted['Momentum_5'] = filtered_df_sorted.
    ↪groupby('Ticker')['Close'].transform(lambda x: x - x.shift(5))
filtered_df_sorted['Momentum_10'] = filtered_df_sorted.
    ↪groupby('Ticker')['Close'].transform(lambda x: x - x.shift(10))
filtered_df_sorted = filtered_df_sorted.dropna()
```

4.1.3 MODELLING AND EVALUATION

I trained a Support Vector Classifier (SVC) with an RBF kernel, tuning the parameters c and γ using cross-validation. I evaluated my model using accuracy, precision, recall, F1-score, confusion matrix and ROC-AUC as the evaluation metrics.

```
[140]: from sklearn.model_selection import train_test_split, GridSearchCV, \
        ↪StratifiedKFold
from sklearn.svm import SVC
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score, confusion_matrix,
    roc_auc_score, roc_curve, classification_report
)
X = filtered_df_sorted.drop('label', axis=1).select_dtypes(include=[np.number])
y = filtered_df_sorted['label']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, \
        ↪stratify=y, random_state=42)
X_train_svm = X_train.mask(np.abs(X_train) > 1e10, np.nan).dropna()
y_train_svm = y_train.loc[X_train_svm.index]
X_test_svm = X_test.mask(np.abs(X_test) > 1e10, np.nan).dropna()
y_test_svm = y_test.loc[X_test_svm.index]
scaler_svm = StandardScaler()
X_train_svm_scaled = scaler_svm.fit_transform(X_train_svm)
X_test_svm_scaled = scaler_svm.transform(X_test_svm)
param_grid = {'C': [0.1, 1, 10, 100], 'gamma': [0.001, 0.01, 0.1, 1]}
svc = SVC(kernel='rbf', probability=True, random_state=42)
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
grid_search = GridSearchCV(svc, param_grid, cv=cv, scoring='accuracy', \
        ↪n_jobs=-1, verbose=1)
grid_search.fit(X_train_svm_scaled, y_train_svm)
print("Best parameters:", grid_search.best_params_)
print("Best cross-validation accuracy:", grid_search.best_score_)
best_svc = grid_search.best_estimator_
y_pred = best_svc.predict(X_test_svm_scaled)
y_proba = best_svc.predict_proba(X_test_svm_scaled)[: , 1]
accuracy = accuracy_score(y_test_svm, y_pred)
precision = precision_score(y_test_svm, y_pred, average='binary')
recall = recall_score(y_test_svm, y_pred, average='binary')
f1 = f1_score(y_test_svm, y_pred, average='binary')
conf_matrix = confusion_matrix(y_test_svm, y_pred)
roc_auc = roc_auc_score(y_test_svm, y_proba)
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1-score:", f1)
print("ROC-AUC:", roc_auc)
print("Confusion Matrix:\n", conf_matrix)
print("\nClassification Report:\n", classification_report(y_test_svm, y_pred))
```



```
fpr, tpr, thresholds = roc_curve(y_test_svm, y_proba)
plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, label=f'ROC curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
```

Fitting 5 folds for each of 16 candidates, totalling 80 fits

Best parameters: {'C': 100, 'gamma': 0.001}

Best cross-validation accuracy: 0.9790866155879824

Accuracy: 0.9889341875364007

Precision: 0.9854423292273237

Recall: 0.9932279909706546

F1-score: 0.9893198426082068

ROC-AUC: 0.9995911803614456


Confusion Matrix:

```
[[818  13]
```

```
[  6 880]]
```

Classification Report:

	precision	recall	f1-score	support
0.0	0.99	0.98	0.99	831
1.0	0.99	0.99	0.99	886
accuracy			0.99	1717
macro avg	0.99	0.99	0.99	1717
weighted avg	0.99	0.99	0.99	1717



4.2 PROBABILISTIC FORECASTING WITH BAYESIAN REGRESSION

4.2.1 FEATURE ENGINEERING

I added the same set of features to my dataset that I used to model my SVC, while adding some more of my own.

```
[143]: filtered_df_sorted['EMA_5'] = filtered_df_sorted.groupby('Ticker')['Close'].
        ↪transform(lambda x: x.ewm(span=5, adjust=False).mean())
filtered_df_sorted['EMA_10'] = filtered_df_sorted.groupby('Ticker')['Close'].
        ↪transform(lambda x: x.ewm(span=10, adjust=False).mean())
filtered_df_sorted['Open_Close_Ratio'] = filtered_df_sorted['Open'] /
        ↪filtered_df_sorted['Close']
filtered_df_sorted['High_Low_Ratio'] = filtered_df_sorted['High'] /
        ↪filtered_df_sorted['Low']
filtered_df_sorted['Log_Return'] = np.log(filtered_df_sorted.
        ↪groupby('Ticker')['Close'].shift(-1) / filtered_df_sorted['Close'])
filtered_df_sorted['Close_ZScore_30'] = (
        (filtered_df_sorted['Close'] - filtered_df_sorted.
        ↪groupby('Ticker')['Close'].transform(lambda x: x.rolling(30).mean())) /
```

```

        filtered_df_sorted.groupby('Ticker')['Close'].transform(lambda x: x.
↪rolling(30).std())
    )
    filtered_df_sorted['Volume_x_Volatility_10'] = filtered_df_sorted['Volume'] *
↪filtered_df_sorted['Volatility_10']
    filtered_df_sorted['PriceChange_x_RSI_14'] = filtered_df_sorted['Price Change']
↪* filtered_df_sorted['RSI_14']
    filtered_df_sorted['Rolling_Min_10'] = filtered_df_sorted.
↪groupby('Ticker')['Close'].transform(lambda x: x.rolling(10).min())
    filtered_df_sorted['Rolling_Max_10'] = filtered_df_sorted.
↪groupby('Ticker')['Close'].transform(lambda x: x.rolling(10).max())
    filtered_df_sorted['Cumulative_Return'] = (filtered_df_sorted.
↪groupby('Ticker')['Daily Return'].transform(lambda x: (1 + x/100).cumprod()))
def williams_r(df, window):
    highest_high = df.groupby('Ticker')['High'].transform(lambda x: x.
↪rolling(window).max())
    lowest_low = df.groupby('Ticker')['Low'].transform(lambda x: x.
↪rolling(window).min())
    wr = -100 * (highest_high - df['Close']) / (highest_high - lowest_low)
    return wr
    filtered_df_sorted['WilliamsR_10'] = williams_r(filtered_df_sorted, 10)
def stochastic_k(df, window):
    low_min = df.groupby('Ticker')['Low'].transform(lambda x: x.rolling(window).
↪min())
    high_max = df.groupby('Ticker')['High'].transform(lambda x: x.
↪rolling(window).max())
    return 100 * (df['Close'] - low_min) / (high_max - low_min)
    filtered_df_sorted['Stoch_K_10'] = stochastic_k(filtered_df_sorted, 10)
    filtered_df_sorted['Stoch_D_10'] = filtered_df_sorted.
↪groupby('Ticker')['Stoch_K_10'].transform(lambda x: x.rolling(3).mean())

```

4.2.2 MODELLING

I trained the Bayesian Ridge Regression model using Scikit-learn

```

[144]: from sklearn.impute import SimpleImputer
from sklearn.linear_model import BayesianRidge
from sklearn.preprocessing import StandardScaler
import seaborn as sns
filtered_df_sorted = filtered_df_sorted.dropna(subset=['Close_Log'])
feature_cols = [
    col for col in filtered_df_sorted.columns
    if col not in ['Close_Log', 'label'] and filtered_df_sorted[col].dtype in
↪[float, int]
]
X = filtered_df_sorted[feature_cols]

```

```

y = filtered_df_sorted['Close_Log']
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, shuffle=False
)
X_train_bayes = X_train.replace([np.inf, -np.inf], np.nan).mask(np.abs(X_train) >
    ↪ 1e10, np.nan)
X_test_bayes = X_test.replace([np.inf, -np.inf], np.nan).mask(np.abs(X_test) >
    ↪ 1e10, np.nan)
imputer_bayes = SimpleImputer(strategy='mean')
X_train_bayes = pd.DataFrame(imputer_bayes.fit_transform(X_train_bayes),
    ↪ columns=X_train.columns, index=X_train.index)
X_test_bayes = pd.DataFrame(imputer_bayes.transform(X_test_bayes),
    ↪ columns=X_test.columns, index=X_test.index)
scaler_bayes = StandardScaler()
X_train_bayes_scaled = scaler_bayes.fit_transform(X_train_bayes)
X_test_bayes_scaled = scaler_bayes.transform(X_test_bayes)
bayes_model = BayesianRidge()
bayes_model.fit(X_train_bayes_scaled, y_train)

```

[144]: BayesianRidge()

4.2.3 FORECAST AND EVALUATION

I forecasted the next-day closing price with credible intervals, and did the evaluation using Mean Absolute Error (MAE) and Interval coverage probability (how often the actual value lies in the 95% credible interval) as the evaluation metrics.

```

[145]: y_pred_mean, y_pred_std = bayes_model.predict(X_test_bayes_scaled,
    ↪ return_std=True)
lower_bound = y_pred_mean - 1.96 * y_pred_std
upper_bound = y_pred_mean + 1.96 * y_pred_std
y_pred_price = np.exp(y_pred_mean)
lower_price = np.exp(lower_bound)
upper_price = np.exp(upper_bound)
y_test_price = np.exp(y_test)
mae = mean_absolute_error(y_test_price, y_pred_price)
print("Mean Absolute Error (MAE):", mae)
coverage = np.mean((y_test >= lower_bound) & (y_test <= upper_bound))
print("Interval Coverage Probability (95% CI):", coverage)
residuals = y_test - y_pred_mean
lower_q = np.percentile(residuals, 2.5)
upper_q = np.percentile(residuals, 97.5)
lower_bound_emp = y_pred_mean + lower_q
upper_bound_emp = y_pred_mean + upper_q
coverage_emp = np.mean((y_test >= lower_bound_emp) & (y_test <=
    ↪ upper_bound_emp))
print("Empirically calibrated coverage:", coverage_emp)

```

```


lower_price_emp = np.exp(lower_bound_emp)
upper_price_emp = np.exp(upper_bound_emp)
if isinstance(y_test.index, pd.MultiIndex):
    x_axis = y_test.index.get_level_values('Date')
else:
    x_axis = y_test.index
plt.figure(figsize=(14, 7))
plt.plot(x_axis, y_test_price, label='Actual Close', color='black')
plt.plot(x_axis, y_pred_price, label='Predicted Mean', color='blue')
plt.fill_between(x_axis, lower_price, upper_price, color='orange', alpha=0.3,
    ↳label='95% Credible Interval')
plt.fill_between(x_axis, lower_price_emp, upper_price_emp, color='green',
    ↳alpha=0.2, label='Empirical 95% Interval')
plt.legend()
plt.title('Bayesian Ridge Regression: Forecast vs Actual')
plt.xlabel('Date')
plt.ylabel('Price')
plt.tight_layout()
plt.show()
importance = np.abs(bayes_model.coef_)
feature_importance = pd.Series(importance, index=X_train.columns).
    ↳sort_values(ascending=False)
print("Top 10 features by absolute coefficient value:")
print(feature_importance.head(10))
residuals = y_test - y_pred_mean
plt.figure(figsize=(10, 4))
sns.histplot(residuals, kde=True)
plt.title('Residual Distribution')
plt.show()
jb_stat, jb_p = jarque_bera(residuals)
print(f"Jarque-Bera p-value: {jb_p:.4f}")

```

Mean Absolute Error (MAE): 134.53114075832275

Interval Coverage Probability (95% CI): 0.34110289937464466


Empirically calibrated coverage: 0.9499715747583855



Top 10 features by absolute coefficient value:

SMA_7_Squared	0.774469
30-Day Moving Average	0.491534
SMA_30_Squared	0.294454
Volume_Price_Ratio	0.269047
BB_Lower_5	0.243228
MA_5	0.186121
BB_Mid_5	0.186121
EMA_10	0.175410
BB_Upper_10	0.145344
EMA_5	0.137376

dtype: float64



Jarque-Bera p-value: 0.0000

4.3 TRADING STRATEGY (HYBRID MODEL)

Using a Hybrid of the two models, with Entry Rule as buy if SVM predicts up and Bayesian mean forecast > current price, and Exit Rule as addition of stop-loss (e.g., 3%) and take-profit levels, I backtested my dataset, using the same logic from before, reporting the profit and loss, Sharpe Ratio and Drawdown.

```
[150]: svm_index = X_test_svm.index
        bayes_index = X_test_bayes.index
        final_index = svm_index.intersection(bayes_index)
        y_pred_series = pd.Series(y_pred, index=svm_index)
        y_pred_mean_series = pd.Series(y_pred_mean, index=bayes_index)
        lower_bound_emp_series = pd.Series(lower_bound_emp, index=bayes_index)
        upper_bound_emp_series = pd.Series(upper_bound_emp, index=bayes_index)
        backtest_df = X_test.loc[final_index].copy()
        backtest_df['Open'] = filtered_df_sorted.loc[final_index, 'Open'].values
        backtest_df['Close'] = filtered_df_sorted.loc[final_index, 'Close'].values
        backtest_df['SVM_Pred'] = y_pred_series.loc[final_index].values
        backtest_df['Bayes_Mean'] = np.exp(y_pred_mean_series.loc[final_index].values)
        backtest_df['Lower_Bound_Emp'] = np.exp(lower_bound_emp_series.loc[final_index].
        ↪values)
        backtest_df['Upper_Bound_Emp'] = np.exp(upper_bound_emp_series.loc[final_index].
        ↪values)
        backtest_df = backtest_df.reset_index(drop=True)
```

```
[159]: def backtest_hybrid_strategy(
        df,
        stop_loss_pct=0.03,
        take_profit_pct=0.05,
        position_size=0.03,
        transaction_cost=0.005,
        slippage=0.001,
        cooldown_period=3,
        min_svm_prob=0.6,
        svm_proba_col='SVM_Proba'
    ):
        pv = 1.0
        shares = 0
        in_market = False
        entry_price = 0
        cooldown = 0
        portfolio_values = [pv]
        daily_returns = []
        max_drawdown = 0
```

```

peak = pv
for i in range(len(df)):
    if pv <= 0:
        break
    row = df.iloc[i]
    open_price = row['Open']
    close_price = row['Close']
    svm_pred = row['SVM_Pred']
    svm_proba = row.get(svm_proba_col, 1.0)
    bayes_mean = row['Bayes_Mean']
    current_price = close_price
    if cooldown > 0:
        cooldown -= 1
    if not in_market and cooldown == 0 and svm_pred == 1 and svm_proba >= min_svm_prob and bayes_mean > current_price:
        trade_capital = pv * position_size
        if trade_capital > 0:
            effective_entry_price = open_price * (1 + slippage)
            shares = (trade_capital * (1 - transaction_cost)) / effective_entry_price
            entry_price = effective_entry_price
            in_market = True
            pv -= trade_capital * transaction_cost
        if in_market:
            effective_exit_price = close_price * (1 - slippage)
            stop_loss_triggered = effective_exit_price <= entry_price * (1 - stop_loss_pct)
            take_profit_triggered = effective_exit_price >= entry_price * (1 + take_profit_pct)
            exit_trade = stop_loss_triggered or take_profit_triggered
            if exit_trade:
                exit_value = shares * effective_exit_price * (1 - transaction_cost)
                pv += exit_value
                shares = 0
                in_market = False
                pv -= exit_value * transaction_cost
                cooldown = cooldown_period
    pv_end = pv + (shares * close_price if in_market else 0)
    daily_return = (pv_end - portfolio_values[-1]) / portfolio_values[-1]
    daily_returns.append(daily_return)
    portfolio_values.append(pv_end)
    peak = max(peak, pv_end)
    drawdown = (peak - pv_end) / peak if peak > 0 else 0
    max_drawdown = max(max_drawdown, drawdown)
    if not in_market:
        shares = 0

```



```

        if pv <= 0:
            break
        pnl = portfolio_values[-1] - 1.0
        daily_returns = np.array(daily_returns)
        sharpe_ratio = (daily_returns.mean() / daily_returns.std()) * np.sqrt(252)
        if daily_returns.std() > 0 else np.nan
        max_drawdown = -max_drawdown
        return {
            'P&L': pnl,
            'Sharpe Ratio': sharpe_ratio,
            'Max Drawdown': max_drawdown,
            'Portfolio Values': portfolio_values,
            'Daily Returns': daily_returns
        }

```

```

[162]: results = backtest_hybrid_strategy(
    backtest_df,
    stop_loss_pct=0.03,
    take_profit_pct=0.05,
    position_size=0.01,
    transaction_cost=0.01,
    slippage=0.002,
    cooldown_period=7,
    min_svm_prob=0.7
)
print(f"P&L: {results['P&L']:.4f}")
print(f"Sharpe Ratio: {results['Sharpe Ratio']:.4f}")
print(f"Max Drawdown: {results['Max Drawdown']:.4f}")
plt.figure(figsize=(12, 6))
plt.plot(results['Portfolio Values'], label='Hybrid Strategy')
plt.xlabel('Time')
plt.ylabel('Portfolio Value')
plt.title('Portfolio Value Over Time')
plt.legend()
plt.grid(True)
plt.show()

```

P&L: 1.4999

Sharpe Ratio: 3.9390

Max Drawdown: -0.0094

